

UNIVERSITY OF CAMBRIDGE

MASTER'S THESIS

User Authentication for Pico

Author:

Cristian M. Toader

Supervisor:

Doctor Frank Stajano

*A thesis submitted in fulfilment of the requirements
for the degree of 'Master of Philosophy'*

in the

Computer Security Group
Computing Department

May 2014

Abstract

The abstract needs to be written at the end.

Acknowledgements

The acknowledgements and the people to thank go here, don't forget to include your project advisor...

Contents

Abstract	i
Acknowledgements	ii
Contents	iii
List of Figures	v
List of Tables	vi
Abbreviations	vii
Symbols	viii
1 Introduction	1
1.1 Contribution	2
1.2 Prerequisites	3
2 Pico: no more passwords!	4
3 Assessment framework	7
3.1 Framework for evaluating web authentication schemes	7
3.2 Token based authentication framework	8
4 Design	13
4.1 Proposed design	13
4.2 Framework evaluation	16
4.3 Conceptual design threat Model	19
4.3.1 Unified device	20
4.3.2 Dedicated device with shared components	21
4.3.3 Insecure communication with Pico	23
4.3.4 Shared device with shared components	24
4.3.5 Proposed secure implementation	24
5 Implementation Prototype	26
5.1 Authenticator design	26
5.2 Implementation details	27
5.2.1 Main application and services	27

5.2.2	Authentication mechanisms	28
5.2.2.1	Dummy mechanism	33
5.2.2.2	Voice recognition	34
5.2.2.3	Face recognition	36
5.2.3	Owner configuration	37
5.3	Threat model	38
5.3.1	Literature review: Android security	38
5.3.2	Future work	44
5.4	Results	45

A	Appendix Title Here	46
----------	----------------------------	-----------

List of Figures

List of Tables

Abbreviations

LAH List Abbreviations **Here**

Symbols

a	distance	m
P	power	W (Js^{-1})
ω	angular frequency	rads^{-1}

Chapter 1

Introduction

As shown by papers such as [1] written by Adams & Sasse passwords have become increasingly difficult to manage. They are meant to authenticate an user based on a shared known secret. Although in theory as well as in the past this has worked well, in practice they are no longer a viable scalable solution.

Increasing computation power has made passwords easier to brute force. In order to avoid this, authentication systems started enforcing rules such as a minimum number of characters, having at least 1 numeric character, or at least one special character. Furthermore the number of accounts which require passwords has significantly increased since they first became widely adopted. In order to maintain security and not have the exposure of a password affect multiple accounts, each password should be unique. Furthermore, security experts recommend (and system admins sometimes enforce) that passwords are to be changed regularly with something that is not too similar to the previous one.

All these security restrictions and recommendations are in place in order to make the mechanism theoretically secure. There is no focus on the user which is meant to memorise numerous unique complex passwords. As shown in paper [2] most of the times users ignore recommendations that are not enforced by the authentication system or even worse, they write the passwords down and compromise security for usability.

The Pico project [3] by Stajano was designed to make passwords obsolete. It is a hardware token which keeps and generates user authentication credentials. This unburdens the user from the dreaded “something you know” transforming it into “something you have”.

This makes the solution scalable with the number of accounts, as well as secure, since the user may no longer choose “weak” passwords.

What differentiates Pico from other token based authentication mechanisms is its locking and unlocking mechanism. Without introducing the need for a PIN or any other known secret, Pico only becomes available in the presence of its owner. This adds security to the threat model when the device is not in the possession of its rightful owner. Currently Pico relies on the concept of Picosiblings [1] which are small in the presence of which Pico becomes unlocked.

Although Picosiblings are a sensible solution to unlocking Pico, they purely based on proximity to the device. The purpose of this project is to create a new way of authenticating the user to the Pico device. Pico is defined as a device which unlocks and locks automatically only in the presence of the owner. This is a relatively open concept. How would the device detect presence? The following chapters will reveal the proposed solution.

1.1 Contribution

The project work is strongly related to the Pico project designed by Stajano [2]. The following contributions have been made in order to provide an alternative to the Picosiblings user authentication:

- We created a framework for assessing token based authentication mechanisms.
- We designed a new way of authenticating the Pico to the user and compare the proposed solution with alternatives.
- We developed an Android prototype which functions as a platform which allows further contributions.
- We analyse the Pico claims according to the framework developed by Bonneau et al [3] and determine the impact of the solution on the Pico.

1.2 Prerequisites

Although the focus of the project is creating an alternative concept for the Picosiblings solution, the prototype was developed using the Android Java SDK. Additional knowledge regarding biometric authentication mechanisms and signal processing would also be useful but implementation details are outside the scope of this project.

Chapter 2

Pico: no more passwords!

The project has strong connections with the Pico project. Therefore, the following sections aim to go into brief detail as to what Pico is, how it works, and what its properties are. T

In the paper “Pico: no more passwords!” [?], Frank Stajano describes an alternative design to passwords. It is based on a hardware token called “Pico”, which replaces all instances where passwords could be used. The reason for designing this alternative authentication mechanism is that passwords are not as usable and secure as they used to be.

Computation power has increased significantly since passwords became popular. This makes passwords easier to break using either brute-force or dictionary attacks. As a result, restrictions were imposed to make passwords stronger, such as requiring at least an upper case letter or a numeric character. Furthermore, in order to reduce the threat of dictionary attacks, authentications systems recommend, if not enforce, passwords to be harder to guess, meaning they should not contain dictionary words. Furthermore, after imposing all these restrictions, it is also recommended to change them regularly. This rule is sometimes enforced in the case of larger companies.

The changes previously described are reasonable from a security point of view, but completely lack user focus. With an increasing number accounts requiring authentication, whether web based or local, passwords become increasingly difficult to remember. Reusing them is not recommended, as compromising one account would lead to multi-account exposure. This leads to having users remember a large number of unique

passwords, which is not realistic in practice. Therefore security compromises are made by the user such as writing passwords down, or reusing them, therefore breaking the theoretical security of the password.

The token based alternative designed by Frank Stajano improves the security offered by passwords. It transforms something you know in something you have. This makes the authentication memory effortless. Unlike other hardware tokens, Pico is theft resistant thanks to small devices called Picosiblings. The Pico uses wireless communication with the Picosiblings and becomes unlocked only in their presence. It is also responsible for generating authentication credentials therefore ensuring that they are not weak or reused.

The Pico device uses for authentication purposes two buttons and a camera. The buttons are used for pairing and creating a new account. The authentication process is multi channel. It relies on a visual code created by the app in order to transmit its public key and unique id. The rest of the authentication process is performed over radio communication via nonce challenges to prove that the other participant has the private key corresponding to the public. The authentication protocol is described in detail in paper [?].

Authentication between the Pico device and its owner is performed using small devices called Picosiblings. These are meant to be embedded in everyday items that the owner carries around, such as earrings, rings, keys, chains, etc. Communication with the Picosiblings uses short range radio. The initialisation protocol for Picosiblings is called the resurecting ducklings [?]. Further enquiries are performed using picosibling pings once every predefined time interval. This tests whether the picosiblings are within the range of the Pico.

The communication between Pico and its Picosiblings is meant to reconstruct the Pico master key. Each Picosibling has a k-out-of-n secret used by the Pico device to reconstruct its master key. Pico tracks each Picosibling based on a timer. If the timer expires and the Picosibling is unable to generate a response, the secret is deleted.

There are two special shares with a longer time out period: a biometric authentication mechanism, and a remote server communication. These are used for special cases such as having your Picosiblings and Pico both stolen. The biometric authentication would only

allow the new user to have limited access time, while the remote server communication can be used to lock Pico remotely.

Chapter 3

Assessment framework

The purpose of this chapter is to create a framework for assessing token based authentication mechanisms. Using this framework we can then compare existing solutions. Having the framework as a performance compass we can then continue by designing an alternative to the Picosiblings unlocking mechanism.

3.1 Framework for evaluating web authentication schemes

In the paper “The Quest to Replace Passwords: A Framework for Comparative Evaluation of Web Authentication Schemes” [?] the authors develop a framework for evaluating web based authentication mechanisms. The purpose of the framework is to identify authentication schemes which outperform passwords. The framework is intended to provide a benchmark for future web authentication proposals.

The framework focuses on three classes of properties which are abbreviated as UDS: usability, deployability, and security. Each class contains a set of properties, totalling a number of 25 benefits. A mechanism may either offer, quasi-offer, or not offer a property. Properties which are not applicable to a mechanism are marked as “offered” to simplify the framework.

Using the framework to evaluate 35 password replacement schemes shows that no scheme is dominant over passwords. According to the evaluation, passwords score perfectly in deployability. They score reasonably in terms of usability, excelling in properties such

as: nothing-to-carry, efficient-to-use, and easy-recovery-from-loss. In terms of security however, passwords don't perform as well, only receiving points in resilience-to-theft (not applicable), no-trusted-third-party, requiring-explicit-consent, and unlinkable. The full list of properties and their description can be found within the paper itself.

Biometric mechanisms receive mixed scores on usability. None of them have the infrequent-errors property which is a precision problem related to false negatives. More importantly if the biometric data is exposed by malware for instance, the authentication mechanism may not be used by the user any more. They score poorly in deployability due to the additional hardware required. In terms of security they perform worse than passwords. Replay attacks can be used by an attacker using a recording of some sort in order to trick the sensor. They are not resilient to theft, since they require an additional device. The fact that they uniquely link the owner to the recording means that the owner may be linked back to the data, therefore not granting the "unlinkable" property.

The paper notes that the memory-effortless property versus nothing-to-carry is only achieved by biometric schemes. None of the mechanisms manage to fully achieve memory-effortless and be resilient-to-theft. This is due to the fact that most mechanisms replace something you know with something you have.

The authors do not produce aggregate scores or rankings. This is due to the fact that not all properties are equal in importance, but different properties would have different weights depending on the scheme's application domain.

Combining schemes is mentioned as a two factor arrangement. This would result in a mechanism which in terms of usability and deployability would only have the properties which are granted by both schemes. In security however it would have the properties of both mechanisms. As shown in the paper, according to [?] the presence of a two factor authentication would lead the user to creating weaker passwords.

3.2 Token based authentication framework

Web based authentication mechanism is initiated locally and performed locally. In contrast, token based authentication is initiated and performed locally, leaving no room for man in the middle attacks or any other 3rd party participation. For this reason, a

subset of the properties described in paper [?] by Bonneau et al should also be present in the framework we have developed.

The following list shows what properties of the framework developed by Bonneau et al are relevant to token based authentication mechanisms:

Memory-effortless

Different types of tokens would have different results. For instance the RSA SecurID [1] token doesn't require any authentication, while the FIDO (Fast IDentity Online) Alliance [2] may use a PIN requiring a known secret.

Easy-to-learn

Token authentication mechanisms may have different learning curves. As an example a CAP reader is fairly easy to use, while a Pico device may prove more difficult for the inexperienced.

Efficient-to-use

Time required by the token user authentication mechanism may differ from one type of authentication to the other. The time required for registering a new user or unlocking the token for its owner should be reasonably short.

Infrequent-errors

The token unlocking mechanism may reject true positives. If the number of false negatives is reasonably low, then the mechanism has this property.

Easy-recovery-from-loss

The user's ability to get another token which uses the same authentication mechanism. Tokens which unlock using biometrics for instance, if not properly secured may lead to the user's inability to use that mechanism again.

Accessible

The ability for all users to use the authentication mechanism. As an example, PINs may be entered by any user regardless of disabilities, on the other hand other biometric mechanisms may not be available.

Negligible-cost-per-user

The total cost enquired by the user in order to use the authentication mechanism.

Mature

It refers to the number of users that have successfully used the mechanism, open source projects based on the mechanism, and any other usage by a third party which did not participate in the development of the scheme.

Non-proprietary

Anyone can implement the token unlocking scheme without having to pay royalties to anyone else.

Resilient-to-physical-observation

An attacker should not be able to impersonate the user after observing them authenticate.

Resilient-to-targeted-impersonation

An attacker should not be able to impersonate the user using knowledge about the user, or previous recordings of his biometrics.

Resilient-to-throttled-guessing

The resilience to an attacker automating a guessing process in order to brute force the unlock of the token.

Resilient-to-unthrottled-guessing

An attacker which only physical access to the token cannot guess the required unlocking resource.

Resilient-to-internal-observation

An attacker cannot tamper with the token in order to intercept user input. Furthermore it is impossible for the attacker to gather the input from within the token's storage.

Requiring-explicit-consent

The authentication mechanism requires explicit consent from the user in order to become unlocked.

Unlinkable

The unlocking mechanism does not generate data which if leaked would compromise the identity of the user. ...

The properties described above are derived from the original framework presented by [?]. Additional details relevant to each property, including when a property is only quasi (partially) satisfied may be found in the original work by Bonneau et al. Some properties from the original work, such as Nothing-to-carry or Server-compatible, do not apply for token unlocking schemes and therefore are not included.

Although the framework by Bonneau et al provides a good base set of properties, a few others are needed in order to fully characterise token unlocking mechanisms. The following list is part of the project's contributions to the overall evaluation framework.

Continuous-authentication

The concept, although mentioned, is omitted from the framework developed by Bonneau et al [?]. It stressed a bit more as part of the benefits of Pico by Stajano [?]. This is a property belonging to the security category of the original framework. The property is satisfied if, once authenticated, the user remains authenticated to the token for as long as he is in its presence. This is similar to an authentication session with the added property that the session remains active for as long as the user requires it. The property should be satisfied by mechanisms which may re-perform authentication in a non explicit way, leaving the user unaware of the underlying process.

Multi-purpose-unlocking

This is a security category property. If satisfied, the unlocking mechanism may allow for multiple types of unlocking based on the user confidence or identity. This is something that may be characteristic for mechanisms which involve biometrics or accounts with multiple security levels.

Availability

This is an usability property. If satisfied, the user has the ability to use the unlocking mechanism in any circumstance. As an example gait recognition would only function while moving on foot, or a recognizer that restricts access based on pulse would not satisfy this requirement. A mechanism requiring only a PIN on the other hand would work in any circumstance.

In order to demonstrate how the framework works we have assessed a number of token based authentication mechanisms, including the original Picosiblings solution. Results are shown in figure ??.

At the end of this piece of work a new framework for the evaluation of token unlocking mechanisms was developed. Existing properties have been identified from the literature and added to the framework together with original work. An initial evaluation was made for existing token unlocking mechanisms which will serve as a benchmark for the proposed solution.

Chapter 4

Design

4.1 Proposed design

The design proposed by Stajano [?] requires Pico's locking and unlocking mechanism to be performed without the use of any password mechanisms. Furthermore support needs to be added for continuous authentication. What this means is that Pico needs to detect the presence of the user in a non obtrusive way.

Given these properties, my idea is to combine multiple continuous authentication mechanisms in order to compute a confidence level. If the confidence level is satisfactory then the Pico is notified that the owner is present and becomes unlocked. This would provide a non-obtrusive continuous authentication mechanism that would suit the Pico requirements for satisfying its claims.

Since the process needs to be continuous and non-obtrusive, most of the mechanisms rely on biometric data such as iris recognition, face recognition, voice recognition, gait recognition, and others. Another source of data would be the user's location compared with previous history, and other patterns that give with a certain confidence whether the owner is in the presence of the Pico.

The approach offered by this project is different from simply stating that Pico is using biometric data as an unlock mechanism. The novelty in this design is based on how the data is combined in order to compute the owner confidence level. Each is assigned a different weight based on the level of trust it offers in identifying the user. This doesn't

necessarily need to be the precision of the mechanism but it would be a good indicator for choosing a value.

Just as the Picosiblings solution, the Pico may have a wide range of supported biometric inputs. However, not all data is always available or relevant. As an example gait recognition would only work while the user is travelling on foot. Other biometrics such as iris recognition may not always be available based on how the sensors are integrated and carried by the user. For this reason all mechanisms will have a decaying confidence level which decreases in time from the last successful recording.

Let us take for example voice recognition which would be sampled every minute. The current weight of the mechanism is 0 so its output is completely ignored. The next sample is recorded showing a confidence of 70% that the owner is present. Upon this recording, the mechanism weight is updated to it's original value. For the next 10 minutes the owner will be silently reading. Since the mechanism does not manage to identify any voice present, the weight of the mechanism decays. The overall result will be impacted less by the voice recognition mechanism up to the point that the recording will be so old, it will no longer be taken into account, or a new successful identification will be performed.

Each mechanism outputs a probability that the recorded data represents the owner of the token. Upon each recording, this probability is updated using Bayes' Law. This process is also known as a Bayesian update. The equation is described below:

$$P(H|E) = \frac{P(H) * P(E|H)}{P(E)}$$

In the equation above:

- $P(H|E)$ represents the probability of hypothesis H after observing evidence E ; this is also known as the posterior probability.
- $P(H)$ represents the probability of hypothesis H before observing evidence E ; this is also known as the prior probability.
- $P(E|H)$ represents the probability that the current evidence belongs to hypothesis H .

- $P(E)$ is the model evidence, and has a constant value for all hypothesis.

We will use the “Law of total probability” in order to compute the value of $P(E)$ in order to accurately compute the posterior probability. The formula is the following:

$$P(E) = \sum_n P(E|H_n) * P(H_n)$$

Using this the Bayes’ Law equation becomes:

$$P(H|E) = \frac{P(H) * P(E|H)}{\sum_n P(E|H_n) * P(H_n)}$$

Our model however, contains only two hypothesis: the recording of the data either belongs to the user, or not. We can therefore consider $P(H)$ to be the hypothesis that the data belongs to the owner and $P(\neg H)$ that the data belongs to someone else. Obviously the value of $P(\neg H)$ is $1 - P(H)$ and $P(E|\neg H) = 1 - P(E|H)$ Using this information, the rule for updating mechanism probability that the recording belongs to the user becomes:

$$P(H|E) = \frac{P(H) * P(E|H)}{P(H) * P(E|H) + P(\neg H) * P(E|\neg H)}$$

Now that we have shown how mechanism probability is calculated, and know that each mechanism has a decaying weight based on the last recording time we can continue to calculate the overall confidence of the Pico. This is performed quite trivially using a weighted sum. The following equation shows the process:

$$P_{Total} = \frac{\sum_{i=1}^n (w_i * P_i(H|E_i))}{\sum_{i=1}^n w_i}$$

The result is then compared with the minimum threshold required by Pico. If the requirement is satisfied, the user is granted access for the current app authentication. Due to the continuous authentication property, the Pico token will continue to ask its authenticator whether the confidence level is still satisfied. Based on the decay rate of the weights and the input data available of the authenticator’s mechanisms this will constantly be recalculated.

At some point the confidence level required by Pico might be too high for the authenticator to grant access. As an example the owner will want to access it's bank account after being silent in a dark room for the past hour. Let us say this would require a confidence level of 95%, while the authenticator may only output a 20% confidence that the user is still present. Given the circumstances, an explicit authentication mechanism may be required from the user in order to increase the current confidence level.

Combining explicit authentication with the current design can be performed consistently with the continuous authentication mechanisms. Whenever an explicit authentication is required, the only difference will be the fact that the user becomes aware of the authentication process. They are prompted to pass an authentication challenge (i.e. facial recognition, voice recognition). This would guarantee valid input for the authenticator which may then proceed to compute an accurate score.

4.2 Framework evaluation

We will continue by evaluating the new proposed scheme with the token unlocking framework defined in the previous chapter.

Memory-effortless: Quasi-satisfied

TODO: ask about this!

Easy-to-learn: Satisfied

In order to satisfy Pico's property of continuous authentication, all mechanisms part of the scheme I developed also need to have this property. Therefore the authentication process is non-transparent to the user, and therefore there is nothing to learn.

Efficient-to-use: Satisfied

The authentication data is collected either at fixed time intervals, or is fired during special events. The authentication process however, does not fully depend on recent data. A response may be generated without any recent authentication data. Therefore the time spent by the mechanism to generate a response is immediate.

Infrequent-errors: Quasi-satisfied

Given that the scheme depends on biometric mechanisms, the quality of the errors

is as good as the underlying biometrics. If the scheme cannot generate a high enough confidence an explicit biometric challenge will be issued for the user to satisfy. Since the original biometric mechanisms do not have this property, to some extent neither will the scheme I have designed. However, the scheme is combining multiple biometrics results with different score weights based on importance and accuracy. This is much more likely to be accurate, which is why I will mark this as Quasi-satisfied. For a more accurate response, the design needs testing with a high quality prototype.

Easy-recovery-from-loss: Not-satisfied

Token based mechanisms in general do not have this property due to the inconvenience of replacing the token. In our case, the property is also not satisfied. The user would have to re-acquire a new token and reconfigure the owner's biometric data. Furthermore based on the mechanism, such as location settings or gait recognition, the token is likely to require an adaptation period.

Accessible: Satisfied

Due to the fact that the scheme is based on multiple biometrics and location settings, I consider this property to be Satisfied or as a very least Quasi-satisfied. The scheme functions based on available biometrics, without having any predefined solutions. It is highly unlikely that the owner cannot generate any of the available biometric inputs, especially for some such as "face recognition".

Negligible-cost-per-user: Quasi-satisfied

This property depends on the way in which the scheme is implemented. If the implementation is based on high quality sensors embedded in items of clothing and such, then the property is not satisfied. If the implementation reuses sensors that the user already possesses, the the property is fully satisfied as the cost is 0. An example of such an implementation would be an Android application/service possibly using the future Google Glass hardware.

Mature: Not satisfied

This property is not satisfied as the project is at the level of a work in progress prototype. The design is quite fresh and was not implemented by any third party. Neither was is reviewed by the open source community or has had any user feedback.

Non-proprietary: Satisfied

Anyone can implement the scheme without any restrictions such as royalty checks or any other sort of payment to anyone else.

Resilient-to-physical-observation: Satisfied

Since the mechanism is based on biometric data, simple observations from an attacker cannot lead to compromising the user's authentication to the token. The attacker would have no way of reproducing the input through simple observation.

Resilient-to-targeted-impersonation: Quasi-satisfied

Saying that the scheme Quasi-satisfies this property is a bit generous. Each of the mechanisms is vulnerable to a replay attack. An attacker may record one of the user's biometric and replay it as a token input. However, given that the token uses multiple mechanisms, some of which being location based, this is a highly unlikely occurrence. The only vulnerable point would be the explicit authentication mechanisms, which carry a lot of weight.

Resilient-to-throttled-guessing: Satisfied

The amount of throttled guessing required for the user to break one of the biometric mechanisms is far too large for this to actually be a threat.

Resilient-to-unthrottled-guessing: Satisfied

Given that the Resilient-to-throttled-guessing property is satisfied, this property is also satisfied.

Resilient-to-internal-observation: Satisfied

This property does not apply to this scheme.

Requiring-explicit-consent: Quasi-satisfied

The user may opt out from the scheme being active. However, in order for the Pico continuous authentication property to be available, non explicit periodic authentication needs to be performed based on biometric data. This should be viewed as the user consenting for authentication when deciding to use the device. Since there is some uncertainty regarding what explicit consent means in this case, only the property is marked as only Quasi-satisfied.

Unlinkable: Not-satisfied

Just as any of the biometric mechanisms, this property is not satisfied by the mechanism. The authentication data maps uniquely to the owner of the token.

Continuous-authentication: Satisfied

The mechanism was designed with continuous authentication in mind. Data is collected periodically with a confidence weight decaying over time. This allows for the token to be used at any time based on current existing data. The only exception breaking the model would be the explicit authentication mechanisms, but these could only be triggered at the beginning of an authentication process using the token.

Multi-purpose-unlocking: Satisfied

This property is fully satisfied by the authentication mechanism. It allows the token to grant access to different authentication accounts based on the precomputed level of confidence that the owner is present.

Availability: Satisfied

Some mechanisms are not always available even though enabled, especially due to the continuous authentication property. As an example gait recognition while sitting in an office. However, the scheme may use a multitude of mechanisms with the unlikeness that all of them are unavailable. For instance location history may predict with a certain confidence that the owner still in possession of the token. This property is aided by the explicit authentication mechanism which requires explicit input from the user.

Comparing this score with the other token unlocking mechanisms... TODO!

4.3 Conceptual design threat Model

An accurate threat model on the proposed unlock mechanism must start by analysing the set of assumptions made about the mechanism. From there we can identify available threats and how the scheme can be exploited in order to unlock the Pico without owner permission. Throughout the threat model we will explain how relaxing the initial set

of assumptions may change the security outcome. Each model is analysed from an Availability, Integrity, and Confidentiality.

It is important to note that confidentiality is an important category in this evaluation. This is because the device will store sensitive biometric data which is directly linkable to the user. Losing this data, especially in plain-text, would disable the user from ever using the biometric device for which the data was leaked. This is due to the fact that the leaked data could always be replayed, successfully tricking the biometric mechanism.

In each subsection, the model will obviously only introduce issues with the mechanism. Therefore when reading a subsection, the issues are not only those currently presented, but also those from previous subsections that lead up to that point.

4.3.1 Unified device

We will start from the assumption that the unlock mechanism is integrated on the same device with the Pico. The device is assumed to be dedicated and runs no other software. Furthermore, the set of available sensors will also be integrated within the device. Alternatively there may also be peripheral sensors, with no way for an attacker to tamper with the communication to the authenticator.

Availability

From an availability point of view, an outside attacker cannot create a denial of service scenario. Interactions with the device are performed physically, so therefore the device cannot be made unavailable while in the possession of its owner. If the Pico would temporarily lose ownership, from a software perspective it would lock up due to mismatching biometric and location data, but would become available again in the presence of the owner.

Only hardware modification would affect data availability. Simply disconnecting the sensor would not affect the scheme's ability to generate viable results due to the fact that multiple biometrics are used. However an attacker could modify a sensor to output wrong data, tricking it into saying the user is never the owner. This would create a

successful denial of service attack path where a few sensors output that the owner is never present.

Integrity

Communication paths are not accessible from the outside and therefore cannot be tampered with in order to modify data. Furthermore the device is not running any other software and is therefore safe from any malware attacks.

Only physical tampering with the device would change data integrity. Modifying one of the sensor's and changing its output to some random data would be undetectable by the mechanism.

Confidentiality

No software access as well as no communication with the outside (i.e. wired communication) means that data is safe as long as the device is with its owner.

If the device were to be lost, Storage data should be kept encrypted, similar to the way Ironkey [] protects its data. Unfortunately an attack path may already be identified which is due to the fact that using this model the decryption key needs to be stored on the device. An attacker which has hardware access could therefore extract the key and decode the data. The original Picosiblings solution circumvented this approach by keeping

4.3.2 Dedicated device with shared components

We will relax the original set of assumptions by saying that the communication path with the sensors is no longer secure. Furthermore the sensors may be shared with other owners, via a wireless communication link for example. Another feasible scenario is that although sensors are located on the same device as the Pico, the Pico application is fully compartmentalised from the outside world.

What we are trying to stress with this scenario is that the sensors are no longer part of a trusted secure box, but are outside and communication with them, as well as their input may no longer be secure.

Availability

Since the sensors are no longer dedicated, other users may access sensor data. Depending on the hardware and software platform supporting the sensors, this may lead to a denial of service attack on the scheme. For example, if the sensors may only have one owner at a time, an attacker may request data from all sensors keeping them locked from the biometric authentication mechanisms. If the system is built in such a way, then there is nothing the scheme could do to prevent this other than keep the sensors constantly locked for itself. However since the model is built on the concept of shared sensors, this might not be a feasible solution.

Furthermore, communication paths are no longer dedicated. Whether the communication channel is radio or pure software, this introduces a new attack path. A “man in the middle” type of attack may be performed where information data from the sensors is dropped and replaced with bad data. This would create a scenario similar to the one in the previous section, but without the need for physically modifying the sensors.

Integrity

Having shared communication paths with the sensors means that data integrity may be compromised from outside. This goal would be achieved in the previous model only by physically modifying the sensors. Furthermore if the sensors are on the same device as the Pico, malware may modify output data leading to unsuccessful mechanism authentication.

Since Pico and the authenticating mechanism are fully compartmentalised from the outside, their communication is still secure. This compartmentalisation however needs to include all types of storage and communication.

Confidentiality

Unfortunately having shared sensors introduces quite a big confidentiality issue. Given that the sensor data required for authentication is shared, nothing would stop an attacker from collecting just as the Pico unlocking mechanism would. This data could then be replayed to the authenticator in order to unlock the Pico.

This is quite a critical issue. An example of feasible attack pattern would be. A piece of malware analyses when the sensors are locked, and makes assumptions as to when the Pico authenticator is locking them. Based on these assumptions the malware then captures sensor data immediately after the lock was released therefore capturing a possibly valid sample of data.

A more elaborate piece of malware could detect patterns such as time intervals or events that trigger sensor locking. Knowing these patterns it could therefore lock the sensors and gather data just before the Pico authenticator would, and then trick the authenticator by sending it a replay or possibly modified data.

Yet another scenario in these circumstances would be to send the Pico authenticator constant bad data and anticipate the trigger of an explicit authentication request to the user. By locking the sensors at that key time the piece of malware could acquire a high quality data sample. Since most of the mechanisms used by the scheme are biometrics, that data sample would represent permanent damage to the user, as an authentication mechanism using that type of biometric could be replayed in any circumstance.

Since the Pico unlocking mechanism is fully compartmentalised, access its storage is secure and therefore any stored credentials are fully protected.

4.3.3 Insecure communication with Pico

This is a special case model which assumes that Pico and the authenticator we have developed are communicating over an insecure channel. The only thing this case needs to consider is the communication between the two participants.

Availability**Integrity****Confidentiality****4.3.4 Shared device with shared components**

We will relax the model even more in order to better fit reality constraints when implementing the mechanism. In this model, Pico and its authentication mechanism reside in a computing model with shared storage resources. The security of Pico and its authenticator may only be as good as the underlying OS. In order to have a meaningful use-case scenario.

Availability**Integrity****Confidentiality****4.3.5 Proposed secure implementation**

A secure proposed implementation is viable using an Android telephone running a TrustZone enabled ARM processor available in ARMv6KZ [] and later models. This device would essentially be divided into two “worlds”: the normal world running the untrusted Android OS, and the trusted world running a small operating system written for TrustZone. Both operating systems are booted at power up. In addition the TrustZone OS loads a public/private key pair which is inaccessible from Android.

Ideally Pico would be implemented with its authenticator within TrustZone. This would essentially guarantee complete separation from a memory perspective leaving any sort of malware attack impossible via memory.

Persistent memory is however required in order to store data for each individual biometric mechanism used in the authentication scheme. Unfortunately this type of memory is not protected by the TrustZone OS and constitutes a way for a third party to attack

the scheme. However, we could use the TrustZone OS key pair in order to encrypt biometric data on disk. Even though this data is available from Android it would be fully confidential. If properly stored within Android, the OS may even protect its integrity from outside attacks.

Let us consider however that the Android OS has been completely compromised by the attacker and is therefore “hostile”. Under these circumstances data confidentiality can still be fully guaranteed. The TrustZone public key could still be used in order to encrypt the biometric data before writing it to disk. Attacks from a memory perspective may only be performed by modifying data stored on disk. This may only lead to a denial of service for the owner, but not a confidentiality breach.

Let us briefly discuss any issues using the availability-integrity-confidentiality framework.

Availability

Only plausible attacks are denial of service through deleting biometric cache files from disk. This would require constant reconfiguration for the Pico scheme, making the Pico unavailable.

Integrity

Data integrity may only be altered from cache files on disk.

Confidentiality

No known attacks on data confidentiality other than capturing sensor data just as the authenticator would. However this would be possible with or without the Pico being present.

Chapter 5

Implementation Prototype

Thus far we have developed a new Pico authentication scheme and assessed it using our own token unlocking framework. We then have performed a threat model from an availability, integrity, and confidentiality perspective and have suggested the safest implementation which would be as feasible as possible for the user to adopt.

In this chapter we will described design and implementation details for the prototype of the proposed scheme. The implementation platform will be the Android OS, which uses a Java based SDK for application development.

5.1 Authenticator design

The user authenticator for Android is designed to work as a bound service called UAService. Periodically the service outputs to registered Pico clients the status of the authentication process. Any application may be a client as long as it registers with the service. Furthermore, explicit authentication update requests may be performed by the Pico client.

Since different authentication mechanisms require different update periods, we have chosen each mechanism to be represented by an independent service. This allows for more flexibility such as periodic sampling with different intervals. Another feasible use case for example would be performing voice recognition based on the first few seconds of an outgoing or incoming call. This would require a service that is triggered by a PHONE_STATE intent.

Each authentication mechanism service is started and managed by the UAService. Communication between the UAService and each authentication mechanism is enabled through intents. Using this communication link, requests can be made from each individual authentication mechanism in order to get the current confidence level. This value is equal to the probability that the owner is present, multiplied by the weight carried by the mechanism. Given that each mechanism runs as an independent service, weight decay may easily be performed using an AlarmManager or simply a function which is called periodically within the authentication thread.

Either periodically UAService gets the confidence level and weight from each mechanism. It then calculates the overall result. If the result is above the threshold requested by the Pico client, a “Message” is passed back saying that Pico should unlock. Otherwise a negative result is returned, letting the Pico know it should be locked.

5.2 Implementation details

5.2.1 Main application and services

The user authenticator for Android is designed to work as a bound service. According to the Android documentation a bound service exposes functionality to other application components and as well as external applications. It is developed as a regular service which implements the onBind() callback method to return an IBinder. The service lives only as long as a component is bound to it. The service implementation class is called UAService.

The UAService is a central node in the application. It is a bound service for any Pico client which wishes to register for events. Furthermore, it binds any authentication mechanism that is available, enabling it for authentication.

The UAService periodically broadcasts intents to registered clients saying if the Pico should be locked or unlocked. The following interface is exposed to available Pico applications through the “what” parameter of the “Message” class:

MSG_REGISTER_CLIENT

Used for registering a client. The “Message” should have as the “arg1” parameter

the level of confidence required for unlocking. This value should range from 0 to 100. Any values outside these limits will be truncated within the range.

MSG_REGISTER_CLIENT

Used for any application to unregister as a listener for this service. No additional parameters required.

MSG_GET_STATUS

Used by any application when an authentication request is needed. Although the service periodically broadcasts to its registered clients what is the authentication status, explicit requests may also be performed using this “Message”.

UAService interacts with AuthMech objects in order to communicate with an authentication mechanism. Each object is responsible for interfacing the communication with an authentication mechanism. A valid authentication mechanism service needs to extend the AuthMechService abstract class which defines a standard way of communication with the UAService.

Each AuthMechService is programmed as a bound service. UAService binds these services through AuthMech objects. Each AuthMechService exposes the following message passing interface:

AUTH_MECH_REGISTER

Used for registering the UAService service as a client to the AuthMechService.

AUTH_MECH_UNREGISTER

Used for unregistering the UAService service as a client to the AuthMechService.

AUTH_MECH_GET_STATUS

Used by the UAService in order to request the authentication confidence from the AuthMechService. The value will be returned in the `arg1` parameter of the Message passed.

5.2.2 Authentication mechanisms

In order to create a functional prototype, we implemented a couple of mechanisms. The focus of the project is not the quality of the biometric mechanisms involved in the

prototype, their sole purpose being to demonstrate a proof of concept. Android devices offer a wide range of sensor data such as GPS, accelerometer, camera, and microphone.

Based on the sensor data offered by Android devices, a wide range of biometric mechanisms can be developed. A non extensive list may include face recognition, voice recognition, iris scanning, keystroke analysis, gait recognition, and many others. The scheme however, requires a clear predefined list of mechanisms offering continuous authentication as well as explicit.

A number of continuous authentication mechanisms may be developed using solely the standard sensors offered by Android devices. The following non-extensive list was achieved, with details regarding what each mechanism means and how it should be implemented:

Face recognition

This mechanism was also implemented for the purpose of the project, and further details are offered in the following sections. The idea is that based on user behaviour, sampling of the user's face can be performed without any explicit requests. For instance when an user is unlocking the phone it is highly likely that he will be looking at the screen. This creates a good opportunity for the authentication mechanism service to capture an image and determine the confidence level that the unlocker is the actual user.

Voice recognition

This mechanism was also implemented for the purpose of the project, and further details are offered in the following sections. Note that voice sampling does not necessarily imply a voice password of any kind. Voice can be analysed from a feature's perspective, regardless of the words being spoken. Voice sampling can be performed at any time. With a frequent enough sampling rate, the owner of the device is likely to be present in most voice recordings. For even better confidence the mechanism should be implemented to start recording when a call is either made or received. On Android this can be achieved by implementing a listener for a `PHONE_STATE` intent.

Iris scanning

Similar to face recognition, this can be implemented by taking advantage of user

behaviour while using the phone. When the phone is unlocked, the user is very likely to face the front camera, allowing for a good face capture. The only problem with this mechanism is the quality of pictures offered by most phones. If the sampling quality is not sufficiently good, meaningful features from the iris may not be extracted, making the confidence level of the mechanism relatively low.

Keystroke analysis

The principle of keystroke analysis is based on the patterns in which the user types on his mobile phone. Different features can be extracted here, such as: the amount of time the user takes to type letter sequences, words, or individual letters, words per minute, frequent used words, and many others. Based on these features a confidence level can be generated (not carrying a considerable amount of weight). This is harder to implement using solely the Android SDK. A good starting point would be to have a keyboard application developed for the user, which also communicates with the authentication mechanism. Obviously if the keyboard is disabled by an attacker this should still be considered, especially if the authenticator was originally configured to listen for input.

Gait recognition

This mechanism is based on the concept of analysing individual walking patterns. Different people walk in different ways, which even though may not be entirely unique for every individual, would still provide some confidence level regarding the user of the device. In the lack of an existing reliable library, efforts have been made to implement this mechanism, unfortunately unsuccessful. The implementation requires accelerometer data from the device, which needs to be normalised from the sensor's perspective. Android offers activity recognition for walking, driving, or standing still. This is achieved by registering a sensor callback for the `TYPE_STEP_DETECTOR` composite sensor.

Ear shape analysis

Studies have shown [] that the shape of the human ear contains enough unique features in order to perform biometric authentication. Taking advantage on user behaviour when using a phone, accurate images can be captured in order to perform such analyses. Within a few seconds from answering the phone, given no peripherals are attached, the user is going to move the phone towards his ear. Based

solely on timing and/or accelerometer data, accurate pictures could be taken of the user's ear before the camera gets too close. Images captured by such a mechanism could then be used to calculate an accurate confidence whether the owner is the person who is answering the phone.

Service utilisation

This proposed mechanism is not biometric based. It exploits patterns in the Android phone's service and app utilisation. Based on current running applications, services, and the time they were started my create a model where some confidence is given as to whether the owner has changed. This mechanism would only be effective in detecting sudden changes, but may easily be obstructed either by removing the Pico authenticator. Furthermore sudden changes in ownership are not promptly detected which is why the mechanism would have a low weight in the overall scheme.

Proximity devices

A mechanism may be developed which tries to connect with other devices also running the authenticator. The two owners don't necessarily need to know one another for the acknowledgement to be performed. Based on day to day activities, users tend to interact or at least be around a lot of the same people. Whether regular travel schedules, or as a better scenario, working in an office, constantly being in the presence of other known devices should give a confidence as to whether the device is in the presence of the user. This mechanism could only be circumvented by co-workers or friends unlocking the Pico, which is why it should never have sufficient weight to unlock the Pico on its own. In combination with other mechanisms however, it would provide a good sense regarding the owner of the device. If the device is "in good company" there is a good chance the owner is also present. This should be enhanced with time data as to when other trusted devices are recognized. Furthermore, based on the ID of the devices the owner comes in proximity to, the mechanism may have different weights for different devices. As an example, even though travelling with your family on holiday and most of the devices there are unknown, given that a number of frequent IDs are in the proximity of the device, the mechanism should still consider to some extent that it is in the possession of its owner. This would work similarly with the Picosiblings idea,

but each Picosibling is a device running this authentication mechanism which is frequently in the proximity of the owner.

Location data

This mechanism is also non biometric. It is similar to “Proximity devices” and much easier to implement. Based on Android GPS data, the phone may detect whether it is in an usual location or not. Just as “Proximity devices” this mechanism should not carry a high weight in the scheme, especially since it would not provide accurate results in scenarios such as holidays.

Picosiblings The original Picosiblings mechanism may also be used with this scheme. Although not part of the standard set of Android device sensors, if available, a Picosiblings implementation may be included as one of the authentication mechanisms.

Some of the continuous authentication mechanisms may also be used for explicit authentication. Based on the non-extensive list mentioned above, the user may be notified to provide accurate information for the following mechanism: face recognition, voice recognition, iris scanning, keystroke analysis, gait recognition, and ear shape analysis. By notifying the user that he has to provide more accurate authentication data, the mechanisms get a better chance of providing valid results. The decay rate after explicit authentication will be slower in order to maintain the continuous authentication property of the Pico for the duration of the authentication session.

In addition to the mechanisms mentioned above, a number of explicit authentication mechanisms which do not satisfy the continuous authentication property of the Pico may be implemented using the Android SDK. It is important to note that any other mechanisms not included in this list need to satisfy the memory property of the Pico, according to which the user doesn’t need to remember any known secret. A non-extensive list of mechanisms includes the following:

Fingerprint scanner

Devices which may have a fingerprint scanner incorporated, such as the iPhone 5S may use this sensor in order to gather biometric data used for authentication. This mechanism cannot actively be used for continuous authentication due to the fact that the user doesn’t come in contact with the sensors on a regular basis.

A mechanism can therefore request explicit fingerprint data, which would then be compared with the owner's biometric model, outputting a confidence for the authentication. This confidence will be combined in the calculation of the overall scheme confidence just as any other mechanism, the only difference being in terms of weight and decay rate.

Hand writing recognition

The user may be prompted to use the touch screen in order to write a word of his choice. This would guarantee the memoryless property, since the user doesn't need to remember any sort of secret. The handwriting would be analysed with a preconfigured set of handwriting samples in order to determine the confidence level that the owner produced the input.

Lip movement analysis

According to the paper [] by TODO, lip movement during speaking may be used to uniquely identify individuals. Lip movement analysis would be performed similarly as described in the paper. The confidence level that the owner produced the input would then be combined in the authentication scheme. This may also be implemented as a continuous authentication mechanism, with lower success rate expectations due to the way users tend to hold mobile phones, which usually doesn't expose the mouth to the camera.

5.2.2.1 Dummy mechanism

In order to perform tests for different confidence levels, a dummy authentication mechanism was implemented using the `AuthDummyService` class. It extends the `AuthMechService` abstract class, which makes it an independent service just to maintain the application model consistent.

The service contains a data access object (DAO) which in this case only produces random confidence levels within a given range. A thread running within the service makes periodic requests to the DAO in order to mimic an authentication mechanism which periodically samples for data. The service is updated based on the produced value.

When the `UAService` wants to update its overall confidence, it makes a `AUTH_MECH_GET_STATUS` request to the `AuthDummyService` service, which returns the most recent confidence

level multiplied by the current decay factor. The result is combined with the result from the remaining authentication mechanism services.

5.2.2.2 Voice recognition

The voice recognition mechanism is implemented as a `VoiceService` class extending the `AuthMechService` abstract class. When the services `onCreate()` method is called, it starts an authenticator thread which periodically samples data from the device's microphone.

The library used for voice recognition is called `Recognito` developed by Amaury Crickx. It is a text independent speaker recognition library developed in Java. It is by no means one of the best voice recognition libraries, but it was best suited for the purpose of this prototype. The library required minimal additional changes. It claims very good results in scenarios with minimal background noise, such as TED talks [] which it was originally tested on by its author.

In order for the application to compile properly, a subset of the `rt.jar` was required. This is due to “`javax.sound.`” packages included with the library which are not available on Android. Unfortunately “`javax`” is a core library also available in Android, but without any of the sound features. For this reason, although the required packages are included with the application this is purely done to trick the compiler that everything is in place. In reality none of the functionality offered by these packages is used by the application. This is avoided by only using the raw features of the library which require direct sound input without any details regarding sound files and formats.

In order to gather samples compatible with the library and manage them properly, we have created the “`VoiceRecord`” class. This class is responsible of gathering microphone input using a predefined compatible configuration listed in the following listing:

- Sample rate: 44100
- Channel configuration: `AudioFormat.CHANNEL_IN_MONO`
- Audio format: `AudioFormat.ENCODING_PCM_16BIT`

The minimum buffer size required by the class is device dependant and pre-calculated in the constructor. The class wraps an `AudioRecord` object used for gathering microphone

data. Due to the nature of the SDK, the recording is saved as a file which then is loaded into memory whenever needed. Although this is not an efficient approach due explicit loads from disk, it serves the purposes of the prototype.

In order to have a better interface to the Recognito library, a DAO class was created. When initialised, the DAO object loads the owner configuration together with the pre-defined set of background noises. It instantiates a Recognito object and trains it using the loaded data. This is performed using the “createVocalPrint” public function made available by the library.

Using the audio data stored as a “double[]” array and the sampling rate stored in the “VoiceRecord” class, we can then call the “recognize” functionality of the library in order to get the closest match to either the owner, or one of the background noises used for training. The library then returns the closest match given its training data, together with the Euclidean distance to that match.

In order to convert the Euclidean distance to a percentage confidence level, an acceptable Euclidean distance threshold is used. Any result above the threshold is considered too high and is truncated to the threshold level. Using the following formula we then convert the value to a confidence level, which is the equivalent of $P(E|M)$, the probability that the evidence belongs to the model.

$$P(E|H) = 1 - \frac{distance}{THRESHOLD}$$

Dividing the distance over the threshold yields a confidence value between 0 and 1, where 1 is a very large distance and hence a bad result. By using one minus this value we invert the meaning, yielding values between 0 and 1 where 1 corresponds to 100% confidence level.

Having calculated $P(E|H)$ we then proceed by calculating $P(H|E)$ using the formula mentioned in the design section of this dissertation. Whenever calculating the current confidence level, we use the value of $P(H|E)$ multiplied by the current decay rate, a number which is periodically decreased within the service. Due to the message passing mechanism using intents, this value needs to be an integer and is therefore multiplied by 100. The overall result is stored in the service and updated whenever the decaying

weight is modified. When a request is made by UAService, the value is returned using the IBinder message passing mechanism offered by Android.

5.2.2.3 Face recognition

The face recognition mechanism was implemented in the “FaceService” class, which extends the AuthMechService abstract class. It is a service running a thread which periodically collects data from the camera. Each sample is analysed using a face recognition library, and a confidence level is outputted for the current sample. Just as the scheme proposes, this confidence level is multiplied by a weight which is a decaying factor.

The library used for face recognition is a port of the Javafaces library [1]. This was the closest functional library found that was compatible with the Android API. Javafaces is a library written entirely in Java, but which unfortunately makes use of the “javax.imageio.” package which is not available in the standard Android SDK. Since a considerable amount of code needed to be changed, we have created a new library [2] based on the original for the Android OS.

I will briefly present the changes made when porting the Javafaces library. The “BufferedImage” class had to be replaced by its Android counterpart, Bitmap. All BufferedImage references and initialization had to be changed. The API was modified to support direct Bitmap input in order to add more flexibility and lighten the main code of the authenticator. Original data formats for black and white images were assumed to have a single colour channel representing the grey value. This had to be changed within the code in order to reflect the Bitmap convention where all 3 colour channels are present but have the same value. Additional modification were required such as data type mismatches as well as other smaller issues.

Unfortunately, this library combined with the Android SDK does not provide accurate results. This is due to the fact that the library requires a rectangle bitmap perfectly containing the face of an individual. Unfortunately the Android SDK although offers face detection, it only provides the location of the midway coordinate between the eyes, and the distance between the eyes. Using this data alone, an accurate crop cannot be made. As a solution, yet another library would need to be used in order to detect faces and provide more accurate data regarding their location and boundaries.

Every fixed time interval, a thread running within the “FaceService” object samples data from the camera at a fixed time interval. This is performed using an instance of the Camera class. Additional configuration is required based on the orientation of the phone. On the device the prototype was developed [], when the phone is held normally a 90 degree rotation of the image is required.

By default, the Android API does not easily allow for a Camera picture to be taken without any sort of notification to the user. By default both a shutter sound and a visual preview display should be present. The shutter sound can easily be disabled by simply not implementing any shutter callback function. The preview display however proves to be a bit more difficult. The solution used with this prototype was to exploit Android’s option to render the preview image to a “SurfaceTexture” object. This satisfies the API’s request to have a visual display preview for the camera, while the “SurfaceTexture” itself doesn’t need to be displayed on screen. Therefore an picture can be taken from a background service without the user being aware of this event.

A DAO class called “FaceDAO” was developed in order to interface with the Javafaces library port. The authentication thread running within the “FaceService” object periodically captures an image from the camera. The image is then validated using the “FaceDAO” object. The value returned from the Javafaces library is the Euclidean distance to the closest registered user, which in our case is the owner of the device. This distance is handled in exactly the same way as the voice recognition mechanism.

Another problem encountered by face recognition mechanism is the size of the data involved in performing the face recognition. With standard pictures, the application runs out of memory and is closed by the Android OS. In order to fix this issue, Bitmaps collected from the camera are scaled to 50

5.2.3 Owner configuration

In order to configure the biometric authentication mechanisms in a flexible, controlled manner a couple of Android activities were developed. There are used to set the initial owner biometrics based on which the mechanisms will output their confidence levels. These activities use the same DAO classes in order to store the data once it was collected.

Due to the size of the data, which is relatively small, the files can be stored in the application's internal memory, making it inaccessible by other applications.

5.3 Threat model

Even though the scheme implementation is a proof of concept, we will continue by analysing different threat models. This will reveal any flaws behind the concept, allowing for a more robust future implementation.

The purpose of the Pico token is to provide a robust authentication mechanism, without the use of any secrets for the owner to remember. Where the Pico unlocking scheme fits, is correctly identifying the owner of the Pico. Attacks may be performed in the form of malware installed on the device while still within the possession of its owner. The main threat however comes from an attacker having physical access to the user's Pico.

It is important to note that since this is a purely software implementation, physical access may mean either that the attacker is in possession of the phone, or that it may replicate the secrets of the victim's Pico on a separate device. Replicating the Pico secrets would clearly create much more damage for the user from a cost perspective. A total reset of authentication credentials would be necessary for all accounts registered with the Pico device.

5.3.1 Literature review: Android security

In order to perform a valid threat modelling of the scheme, we need to have a better understanding of the Android security model, and the "features" offered by different mechanisms such as Intents and filters based on these intents.

An interesting paper by William Enck et al [1] offers a good description of the Android OS, with a focus on the security aspects of the platform. It is a relatively old paper from 2008 which is the same year of the Android initial release. The initial set of Android open standards was however released sooner, in November 2007, allowing researchers such as William Enck to perform an initial analysis of the system.

Android runs on a port of the Linux kernel. This means that many of the linux security mechanisms such as file permissions, are also part of Android. On top of the kernel lies an

App middleware layer which the Java SDK can use in order to extend the functionality of the hand held device.

Each android application is split into multiple components, with no fixed entry point such as a `main()` routine. Based on its purpose, the SDK defines 4 types of components: activity, service, content provider, and broadcast receivers. More details regarding the purpose of each component can be found on the Android website [1]. Another important architectural note is that components, and even applications, may communicate between each other using Intents. This type of communication is abbreviated as Inter Component Communication (ICC).

An important note in the paper is that when a Service becomes bound by another component it cannot be terminated by an explicit stop action. This provides an useful guarantee regarding the lifetime of a service.

Th two types of security enforcements can be split into two categories: ICC and system level. System level security is mainly defined in terms of the user id (UID) and group id (GID) permissions. According to the Android OS, each is allocated an UID and GID. The security guarantees of this mechanism are the same as those of a linux based system. An interesting example would be a vulnerability of the T-mobile G1 phone browser, which due to this system level security enforcement did not affect any other applications.

The focus of the paper however is ICC security enforcements. These are based on I/O control command of the `"/dev/binder"` special node. Since the node needs to be world readable and writeable, linux cannot mediate ICC. This is performed using labels using a MAC framework enforced by a reference monitor. This defines how apps are allowed to access different components.

Components may be defined as either public or private. This refinement is configured by the `"exported"` field defined in a manifest file. It defines whether or not another application may launch or interact with a component from another application. At the time the paper was written, the `"exported"` field was defaulted to `"true"`. However, as shown in a more recent paper [2] written by Steffen and Mathias in 2013, starting with Linux Android (LA) 4.2 the default of this value was changed to `"false"`, therefore conforming to the `"principle of least privilege"`.

Components listening for Intents need to have a registered filter in the manifest file. This is both convenient and secure for the developers of the system. If however the developer of an application wishes to restrict access to intent objects, the SDK provides user definable Intent permission labels as well as Service hooks. These provide runtime security checks for the application and prevent data leakage through ICC. Using permission labels, the developer may broadcast events to which only components which registered that permission may access. The same principle applies to service hooks, but in this case a component holding some permissions tries to bind on a service which checks these permissions and exposes different APIs based on the permissions held by the binding component.

Another more recent paper was written by Steffen and Mathias [10]. The authors focus on deeper issues of the LA, and how they were solved from one update to the other. However, it is shown that OEMs tend not to update the software of their devices once they have shipped, which poses a number of security issues.

The starting point of understanding Android security and how it is bootstrapped is the five step booting process:

1. Initial bootloader (IBL) is loaded from ROM.
2. IBL checks the signature of the bootloader (BL) and loads it into RAM.
3. BL checks the signature of the linux kernel (LK) and loads it into RAM.
4. LK initialises all existing hardware and starts the linux “init” process [11].
5. The init process reads a configuration file and boots the rest of LA.

The android security model described in this paper [12] is split in two categories: system security, and application security.

The android base system (libraries, app framework, and app runtime) is located in the “system” partition. Although this is writeable only by the root user, a number of exploits which allow root access, such as those shown in [13], have been possible. Android provides a keychain API used for storing sensitive material such as certificates and credentials. These credentials are encrypted using a master key, which is also stored in AES encrypted format. In this case, security needs to begin somewhere, an assumption has to be made

about a state being secure which would allow for multiple security extensions. In this case, the master key is considered to be that point of security. However, given a rooted device, where due to an exploitation a process may be granted root privileges, the master key itself may be retrieved from the system therefore compromising all other credentials.

From the application's perspective, an interesting "feature" which may affect the flow of information within LA is the fact that applications from the same author may share private resources. When installing an application the user needs to accept its predefined set of permissions. Due to resource sharing, a situation may present itself when an application which has permissions for the owner's contacts may communicate with an application which has permissions for internet in order to leak confidential data. A developer may therefore construct pairs of legitimate applications in order to mask a data flow attack.

The LA system offers a number of memory corruption mitigations in order to avoid buffer overflow attacks, or return oriented programming attacks. The following list includes mechanisms enabled in time based on the LA version:

- Implements `mmap_min_addr` which restricts `mmap` memory mapping calls. This prevents NULL pointer related attacks.
- Implements XN (execute never) bit to mark memory as non-executable. The mechanism prevents attackers from executing remote code.
- Address space layout randomisation(ASLR) implemented starting with LA 4.0. This is a first step to preventing return oriented programming attacks. The position of the binary library itself is however static, meaning that after a number of attempts, using trial and error, the attacker may succeed using return oriented programming.
- Position independent and randomised linked (PIE) is implemented starting with LA 4.1 in support of ASLP. This makes position of binary libraries themselves randomised.
- Read only relocation and immediate binding space(RELro) was implemented starting with LA 4.1. It solves an ASLR issue where an attacker could modify the global

offset table (GOT) used when resolving a function from a dynamically linked library. Before this update an attacker may insert his own code to be executed through the GOT table.

A number of security enhancement mechanisms are in place in order to make Android a safer environment for its users. In order to prevent malware within the Android App store (Google Play), a program also known as a “on device Bouncer”. The purpose of the bouncer is to verify apps prior to installation for any malware signature or patterns. Secure USB debugging was introduced starting with LA 4.4.2, which only allows hosts registered with the device to have USB debugging permissions. This mechanism is circumvented if the user does not have a screen lock.

As pointed out by Steffen and Mathias, the Android OS is responsible for 96% of mobile phone malware according to a study from 2012 [1]. The authors claim that this is the case due to 4 big issues of the Android concept:

1. Security updates are delayed or never deployed. This is due to a number of approvals that an update needs to get through in order to be pushed to devices. This implies an additional cost to the manufacturer(OEM) which does not generate any revenue. The majority of teams working on LA are focusing on current releases, and in some cases there is simply not enough time and resources to merge Google security updates to the OEM repository. Furthermore, the consequences of a failed system update for the user may cause problems as serious as “brick”-ing the device which is a huge risk for the manufacturer. All these issues contribute to a severe lack in security updates. Therefore, important updates such as RELro are never pushed to LA 4.0, making the device vulnerable.
2. OEMs weakens the security architecture and configuration of LA by introducing custom modifications before they roll out a device.
3. The Android permission model is defective. As pointed out by Steffen and Mathias, according to Kelley et al [2], most users do not understand the permission dialogue when installing an application. Furthermore, even if they could understand the dialogue, most of the time it is ignored in order to be able to use a new exciting app. According to the same study, most applications are over-privileged due to the developers not understanding what each privilege does. Furthermore, as previously

pointed out, applications from the same owner may share resources, therefore creating a valid data flow attack path.

4. Google Play has a low barrier for malware. A developer distribution agreement (DDA) and a developer program policy (DPP) need to be agreed to and signed by the developer before submitting the application to the Android market. However, Google Play does not check upfront if an application adheres to DDA and DPP. The application is only reviewed if suspect of breaking the agreements. According to [1] there are ways of circumventing the Bouncer program, and as a proof TODO write in their paper [2] a list of applications which are in the Google Play store.

Let us now continue by studying the threat model of the Pico authenticator application. We will consider the security mechanisms presented above as the predefined assumptions made in this model. In order to reduce the threat space we will consider the application is running on a hand held device running Android 4.4.2 with all recent updates.

Availability

Breaking the scheme's availability if the device is in the possession of the attacker is relatively trivial. The application can be uninstalled, or the application data cache can be cleared, therefore removing the owner biometric models for the different mechanisms. In this case the owner is already no longer in possession of their Pico, so basically the Pico is already made unavailable.

Let us continue however and study what can be achieved from a DoS perspective by the attacker from the perspective of the individual user app accounts, which would need to be reset by the owner. In order to gain any sort of access and make credentials reset not possible, or at least have a chance in doing so, the attacker would have to unlock the Pico.

Integrity

Confidentiality

5.3.2 Future work

The application was implemented as a proof of concept. It is developed in order to show that different data may be obtained without the owner's knowledge. Additional improvements can be made in order to increase the confidence level of the authenticator. Furthermore, due to time constraints and unavailability of free to use biometric libraries, a number of mechanisms were not implemented. The list can easily be extended by simply creating a class which extends the "AuthMechService" abstract class.

One way to improve the voice recognition mechanism would be to start sampling data whenever a call is active. This would increase the chances of capturing an accurate sample of the owner's voice. In this context, a better voice recognition library can be used, which supports multiple speakers and/or ignores background noise. If such a library is not available, we can rely on the fact that most of the times people take turns when speaking. For the duration of the call, with a high enough sampling frequency, the individual sampling voice of both participants should be captured. However, it is important to take into account a situation in which the thief is calling the owner on a different phone in order to unlock his Pico.

Immediate improvements can be made to the face recognition mechanism. Just as recommended in the description of the mechanism's implementation, another library which provides more meaningful face coordinates may be used for face detection. Alternatively, and preferably, a different library which performs both face detection and recognition can be integrated with the mechanism.

Another improvement for the face recognition mechanism would be from the data sampling perspective. Instead of capturing images at a fixed interval, pictures should be taken only when the phone unlock event is triggered. While the phone is unlocked it is highly likely that the user will face its front camera. This would provide better chances of processing meaningful data.

5.4 Results

Appendix A

Appendix Title Here

Write your Appendix content here.