

University of Cambridge

Master's Thesis

User Authentication for Pico: When to unlock a security token

Author:

Cristian M. Toader

Supervisor:

Doctor Frank Stajano

*A thesis submitted in fulfillment of the requirements
for the degree of 'Master of Philosophy'*

in the

Computer Security Group
Computer Laboratory

June 2014

Abstract

The abstract needs to be written at the end.

Acknowledgements

The acknowledgements and the people to thank go here, don't forget to include your project advisor...

Contents

Abstract	i
Acknowledgements	ii
Contents	iii
List of Figures	v
List of Tables	vi
Abbreviations	vii
Symbols	viii
1 Introduction	1
2 Pico: no more passwords!	5
3 Assessment framework	9
3.1 UDS assessment framework	9
3.2 Token unlocking framework	11
3.3 Example evaluation	16
3.3.1 Picosiblings	17
3.3.2 PIN	18
3.3.3 Face unlock	19
3.4 Conclusions	20
4 Design	22
4.1 Design requirements	22
4.2 Proposed solution	23
4.2.1 Explicit authentication	24
4.3 Related work	27
4.4 Conclusions	29
5 Implementation Prototype	30
5.0.1 Android development and security	30
5.1 Authenticator design	35
5.2 Implementation details	35

5.2.1	Main application and services	35
5.2.2	Authentication mechanisms	38
5.2.2.1	Dummy mechanism	42
5.2.2.2	Voice recognition	43
5.2.2.3	Face recognition	45
5.2.2.4	Location analysis	46
5.2.3	Owner configuration	47
5.3	Conclusion	47
5.4	Related work	48
Evaluation		50
5.5	Framework evaluation	50
5.6	Conceptual design threat Model	54
5.6.1	Dedicated device with dedicated sensors	54
5.6.2	Dedicated device with shared sensors	55
5.6.3	Insecure communication with Pico	57
5.6.4	Shared device with shared components	58
5.6.5	Proposed secure implementation	58
5.7	Threat model	59
5.7.1	Prototype threat model	60
5.8	Future work	62
 Appendix Title Here		 63
 Bibliography		 64

List of Figures

List of Tables

3.1 Token unlocking framework sample assessment.	21
--	----

Abbreviations

LAH List Abbreviations **Here**

Symbols

a	distance	m
P	power	W (Js^{-1})
ω	angular frequency	rads^{-1}

Chapter 1

Introduction

Passwords are currently the most widely used electronic authentication mechanism. They rely on remembering a secret sequence of characters, and providing it as input for the authentication process. This originally offered a sufficiently secure authentication mechanism. However, as shown by Adams & Sasse [1], the fundamental concept of remembering a secret, makes passwords unsuited for the current technological context.

As Robert Morris [2] emphasises in his paper, there is a constant competition between attackers and security experts. The majority of users try to maximise usability of their passwords by choosing secrets which are easy to remember. These however are susceptible to a number of attacks such as brute force, dictionary, pre-compiled hashes, rainbow tables [3], and others. Security experts were able to slow down the attackers in the past without any impact to the user, but with a constant increase in computational power, passwords became easier to breach.

The main flaw of passwords is that when chosen freely they tend to be short and predictable. In order to maintain acceptable security guarantees, a number of requirements in their creation started to be enforced. Some password mechanisms may require a minimum length, one or more numeric characters, and one or more special characters. Security experts recommend that each account needs to have an unique password. Furthermore, passwords sometimes require to be changed on a regular basis with something not too similar with the original.

As shown by Yan et al [4], without any additional advice to make the password more memorable, users choose weaker passwords. From a theoretical perspective, additional

restrictions would make the mechanism more secure. The users are forced to pick a non-intuitive password, and fully utilise the available character set. This makes dictionary and brute force attacks harder to perform. In practice however, users need to memorise numerous, unique, and complex passwords. As shown by Adams & Sasse [1] maintaining all restrictions and security advices proves not to be feasible, leading to poor practices such as writing the passwords down.

The main problem with passwords is the basic principle of users remembering a secret. If the secret is memorable, than an attacker may brute force it with more ease. If it is too complex, then the user may not remember it. Furthermore, since reusing passwords is not safe and given the memory capacity people have, the solution is not scalable. For all these fundamental reasons, passwords prove not to be a reliable solution for the future and even present.

A large number of alternative to passwords are available. However, as shown by Bonneau et al [5], the main advantage passwords have over other authentication mechanisms are in terms of deployability and usability. A study by Clarke et al [6] shows that although 81% of users agree to an alternative to password based phone unlocking, the majority ignore the existence of available solutions. The main conclusion we may draw is that although passwords are not secure, the cost of replacing them and familiarising with a new authentication mechanism is still too inconvenient.

The Pico project was designed by Frank Stajano [7] with the purpose of replacing password based mechanisms. Pico is a hardware token which generates and manages user authentication credentials. This transforms the problem of knowing a secret into having it. Since anyone in possession of such a hardware token would have access to the owner's accounts, this type of authentication is not very secure. Therefore, Pico adds an additional layer of security by being usable only in the presence of its owner. In a sense a security chain is created where "who you are" unlocks "a secret you have" which is used for authentication.

In order to identify the presence of its owner, Pico communicates with small devices called Picosiblings [8]. These devices are embedded in everyday items that the user carries throughout the day (i.e. keys, necklace, rings). Each Picosibling transmits a secret sequence to the Pico. When all required secrets are gathered, Pico becomes unlocked and may be used by its owner.

Picosiblings are a sensible solution to unlocking Pico. However, they are purely based on proximity to the device. As suggested in the original Pico paper [7] anyone in possession of both Pico and the Picosiblings would have full access to the owner's accounts for a limited amount of time. Some additional security features are included, such as having a remote online server as a Picosibling. However, the main downside of this approach is the fact that Picosiblings do not reflect who the user is, but rather additional things the user has.

The purpose of this project is to design and implement a better token unlocking mechanism for Pico. According to its design, the process should be memoryless, and enable continuous authentication. The token should lock and unlock automatically only in the presence of its owner. The solutions that seem to best fit these requirements are biometric authentication mechanisms. For the purpose of this project we have explored the possibility of combining multiple biometrics and behavioural analysis as part of a unified solution. The output from each mechanism is combined to generate an overall confidence level, reflecting that the owner is still in possession of the Pico.

We will explore and evaluate the original Picosiblings solution as well as other token unlocking schemes. The evaluation will be performed using a framework derived from the work of Bonneau et al [5]. This will enable a formal analysis of the benefits and downsides of the new authentication scheme in comparison with existing mechanisms.

Contribution

In the process of designing and developing a new token unlocking mechanism, a number of contributions have been made. The following list presents a summary of these achievements, with further details in the following chapters.

- We create a framework derived from the work by Bonneau et al [5]. This is used to evaluate a few existing token unlocking mechanisms, including Picosiblings. The data is then used as a benchmark when evaluating the proposed solution.
- We design a new token unlocking mechanism. Although the solution may be used in any type of user authentication, it is presented in the context of unlocking the

Pico token. The design is analysed using the token unlocking evaluation framework. A comparison is made with the original Picosiblings solution. The aim of the dissertation is for the new scheme to achieve better results in at least some categories of the token unlocking framework.

- We develop an Android prototype. The implementation is meant to prove that the design is feasible for implementation using existing technologies. The prototype was not developed for performance purposes. However, power analyses as well as timings of different stages of the scheme were recorded to serve as an approximation of the limitations and downsides of the scheme.
- We analyse and determine the impact of the proposed token unlocking mechanism on the Pico. The analysis is performed based on the original framework by Bonneau et al [5]. One of the proposed goals when designing the solution was to make Pico better in terms of at least one property.

Chapter 2

Pico: no more passwords!

The scope of this dissertation project is to design and implement a new unlocking mechanism for the Pico token, as designed by Frank Stajano [7]. A better understanding of the Pico design is therefore necessary. This chapter aims to go into brief detail as to what Pico is, how it works, and what its properties are.

Pico is an user authentication hardware token, designed with the purpose of fully replacing passwords. Although other replacement mechanisms exist, they are generally focused on web based authentication. The scope of the solution described by Stajano addresses all instances of password authentication, both web based as well as offline.

The motivation behind this project is the fact that passwords are no longer viable in the current technological context. Computing power has grown, making simple passwords easy to break. Longer and more complex passwords are now required. As shown by Adams & Sasse [1], this has a negative impact on the users, which have limited memorising capability.

Another reason why passwords are no longer viable is the fact that they are not a scalable solution. Security experts recommend that passwords should be reused for multiple accounts. However, a large number of computer based services require password authentication. In order to respect security recommendations, users would be forced to remember dozens of unique, complex passwords. A study by Florencio et al [9] performed over half a million users confirms the negative impact of scalability on password quality. Furthermore, passwords are often forgotten or reused across accounts.

When designing the Pico password replacement mechanism, Stajano decides to have a fresh start. He describes that an alternative for passwords needs to be at least memory-less and scalable, without reducing security. In the case of token based authentication, the solution also needs to be loss and theft resistant. The Pico token was therefore designed to satisfy these fundamental properties, as well as other benefits emphasised in a paper by Bonneau et al [5] as well as the original work by Stajano.

As a token authentication mechanism, Pico transforms "something you know" into "something you have". It offers support for thousands of credentials which are kept encrypted on the Pico device. The encryption key is also known as the "Pico Master Key". If the Pico is not in the possession of its owner it becomes locked. In this state, the "Pico Master Key" is unavailable and the user cannot authenticate to any app¹.

Credentials are generated and managed automatically whenever the owner interacts with an app. Therefore, the responsibility of generating a strong and unique credential, as well as memorising it, is shifted from the user to the Pico. No additional effort such as searching or typing credentials is required.

Another important feature offered by Pico is continuous authentication. Traditional password mechanisms provide authentication for an entire session. The user is responsible of managing and closing the session when it is no longer needed. Instead, Pico offers the possibility of periodic re-authentication of its owner using short range radio. If either the Pico or the owner are no longer present, the authentication session is closed.

From a physical perspective, Pico is a small portable dedicated device. Its owner should be carrying it at all times, just as they would with a car key. It contains the following hardware components:

- Main button used for authenticating the owner to the app. This is the equivalent of typing the password.
- Pairing button used for registering a new account with an app.
- Small display used for notifications.
- Short range bidirectional radio interface used as a primary communication channel with the app.

¹For the purpose of brevity, any mechanism requiring user authentication will be called an "app" just as in the original paper by Stajano.

- Camera used for receiving additional data from the app via 2-dimensional visual codes. This serves as a secondary communication channel.

As mentioned before, the Pico main memory is encrypted using the Pico Master Key. It contains thousands of slots used for storing unique credentials used in the authentication process. Each credential consists of public-private key information generated during account creation in a key exchange protocol. The public key belongs to the corresponding app, while the private key was generated when creating the account.

During account creation Pico scans a 2D visual code generated by an app. The image encodes a hash of the app's certificate including the app name and public key. Pico starts the protocol with the app using the radio channel, and the app provides a public key used for communication. The key is validated using the hash from the visual code, and the protocol continues. Pico then initiates a challenge for the app to prove that it is in possession of the corresponding private key, and provides a temporary public key. This protects the identity of the owner, by only showing their public key after the app is authenticated. Only then Pico generates a key pair, sends the public key to the app and stores the key pair.

The account authentication process starts when the user presses the main button and scans the app 2D code. The hash of the app's name and public key are extracted from the 2D image. This information is used to find the corresponding credentials. An ephemeral public key encrypted with the app's public key is sent via the radio channel. The app is authenticated by using this key to require the corresponding (user id, credential) pair. Only after the app is authenticated Pico uses the public key generated during the registration process and authenticates itself to the app.

The locking process is an important aspect of Pico which we have not yet fully described. The token should become unlocked only in the presence of its owner. Currently this is achieved using bidirectional radio communication with small devices called Picosiblings [8]. These are meant to be embedded in everyday items that the owner carries around, such as earrings, rings, keys, chains, etc.

The Pico authentication credentials are encrypted using the Pico Master Key. The key is not available on the Pico and can only be reconstructed using k-out-of-n secret sharing,

as described by Shamir [10]. Except for two shares which will be discussed later, each k-out-of-n share is held by a Picosibling.

Using an initialisation protocol based on the resurrecting duckling [11], each Picosibling is securely paired with Pico. After the initialisation process, Pico sends periodic ping requests to which all registered Picosiblings are expected to respond. During each successful ping, the Picosibling sends its k-out-of-n share back to Pico. If enough secrets are provided the "Pico Master Key" is reconstructed and Pico becomes unlocked.

Internally, Pico keeps a slot array for each paired Picosibling. Each slot contains a countdown value, and the key share provided by the Picosibling. The countdown value is decreased periodically. When it expires, the share becomes deleted. Similarly, if k shares are not acquired before a predefined time-out period, all shares are removed.

Except for the Picosiblings, two additional special shares with a larger time-out period are described by the paper:

- Biometric measurement used for authenticating the owner to the Pico.
- Remote server network connection used for locking the Pico remotely.

The possibility of using a smart phone as a Pico is briefly considered in the paper. This would have the advantage of not requiring any additional devices from the user. Modern smart phones provide all the necessary hardware required by Pico. However, this would be a security trade-off in exchange for usability. Mobile phones are an ecosystem for malware, and they present uncertainty regarding the privacy of encrypted data. This option may still be used as a cheaper alternative to prototype and test, which is something we will make use of in this project.

Chapter 3

Assessment framework

The purpose of this chapter is to create an assessment framework for token unlocking mechanisms. This framework will be used to evaluate existing solutions, including the Picosiblings scheme used by Pico. The analysis of the results can then be used to create an alternative unlocking mechanism for Pico. The project aims to achieve better results in some categories, without necessarily completely outperforming it.

3.1 UDS assessment framework

Similar work to what we are trying to achieve in this chapter was performed by Bonneau et al [5]. The authors create a framework for evaluating web based authentication mechanisms. However, the assessment scheme is not entirely compatible for token unlocking schemes. For example, properties such as "Browser-compatible" do not apply, while others need to be redefined to fit our context. The paper however presents a good starting point for our token unlocking evaluation framework. The remainder of this section will present a brief summary of the paper.

The motivation behind this paper is to gain insight about the difficulty of replacing passwords. An assessment framework is created, and a number of web authentication mechanisms are evaluated. It is a useful tool in identifying key properties of web based authentication schemes. The framework is intended to provide a benchmark for future proposals.

The framework consists of 25 properties divided into three categories: usability, deployability, and security. For this reason, it is abbreviated by the authors as the "\UDS framework". An authentication scheme is evaluated by assessing whether each property is offered or not. In the case where a scheme almost offers a property, the authors mark it as quasi-offered. To simplify the framework, properties which are not applicable are marked as offered.

Since passwords are currently the most widely used authentication mechanism, the results are predictable. Evaluating 35 replacement schemes shows that no scheme completely dominates them. Passwords satisfy all the properties in the deployability category. They score reasonably well in terms of usability, excelling in properties such as: "\nothing-to-carry", "\efficient-to-use", "\and easy-recovery-from-loss". However, from a security perspective passwords don't perform well. They only offer the "\resilience-to-theft"¹, "\no-trusted-third-party", "\requiring-explicit-consent", and "\unlinkable" properties. The full evaluation can be found within the paper itself.

Biometric mechanisms receive mixed scores on usability. None of them offer the "\infrequent-errors" property, due to false negative precision. More importantly if biometric data becomes compromised, the possibility of replay attacks makes the authentication mechanism unreliable. They score poorly in deployability partially because they require additional hardware. In terms of security they perform worse than passwords. Replay attacks can be performed by an attacker using a pre-recording data of the user, making them not "\resilient-to-targeted-impersonation" and not "\resilient-to-theft". There is a one to one correlation between the owner and their biometric recording, therefore the "\unlinkable" property is not offered by these mechanisms.

By analysing the framework results, we see that some authentication schemes, such as security tokens, offer "\memory-efficient" in exchange for the "\nothing-to-carry" property. The only schemes that offer both are biometric mechanisms. This is a consequence of replacing "\something you know" with "\something you are" instead of have. For different reasons no mechanism offers both "\memory-efficient" while being "\resilient-to-theft".

When trying to compute an aggregate score using the framework, not all properties should be equal in importance. Different properties should have different weights depending on the purpose of the assessment. For example, if we would try to find the most

¹Not applicable to passwords

secure authentication mechanism, security properties would have a larger weight in the overall evaluation. For this reason, the authors only provide the means for others to make an evaluation based on their needs. No aggregate scores or rankings are provided in the paper.

The authors mention the possibility of combining schemes as part of a two factor authentication. In terms of deployability and usability, the overall scheme offers a property if it is offered by both authentication mechanisms. In terms of security, only one of the two mechanisms needs to offer the property in order for the two factor combination to offer it as well. However, Wimberly & Liebrock [12] observe that combining passwords with a second authentication mechanism scheme leads to weaker credentials and implicitly less security.

The following section will offer more details on the UDS framework properties which also apply to token unlocking. Further information about the framework are not mentioned in this dissertation for the purpose of brevity. The full list of properties, their description, and the evaluation of a number of mechanisms are provided in the original paper by Bonneau et al.

3.2 Token unlocking framework

Unlike web based authentication mechanisms, token unlocking schemes record and process data locally. For this reason, a subset of the UDS framework properties are also present in the framework we have developed. Those which do not apply, or would apply to every token unlocking mechanism were removed. Some properties needed to be adapted to the new context of a token, and therefore will have a slightly different meaning.

The following list contains all properties of the UDS framework developed by Bonneau et al [5] which we adapted and included in the token unlocking framework. A short description is included for the cases when the property is offered or quasi-offered, as well as a small example.

Memorywise-effortless

Users do not need to remember any type of secret. This includes passwords,

physical signatures, or drawings. The property was originally quasi-ordered if one secret would be used with multiple accounts, but in the case of security tokens this does not apply. As an example the RSA SecurID ² is used in conjunction with a password in order to authenticate the user, and therefore does not order this property.

Nothing-to-carry

The unlocking mechanism does not require any additional hardware except for the token. The property is quasi-satisfied in the case of hardware the user would have carried on a normal basis such as a mobile phone. An example of a mechanism that quasi-orders the property is the Picosiblings scheme which uses small devices embedded in everyday items. Biometric mechanisms that require additional sensors such as a fingerprint reader do not satisfy this property.

Easy-to-learn

Users who use the unlocking mechanism would be able to learn it with ease. The original paper by Bonneau et al [5], assess Pico as not "easy-to-learn" due to the complexity of the Picosiblings management³. PINs or passwords however satisfy this property due to the users' familiarity with this type of authentication.

Efficient-to-use

The amount of time the user needs to wait for the token to be unlocked is reasonably short. This includes the time required to provide the input for the mechanism. The same applies for setting up the token unlocking mechanism, but with a larger time scale. In the case of PINs for example the input and processing time are very low, and therefore the scheme orders the property. Mechanisms based on biometrics however may not, depending on the implementation.

Infrequent-errors

The rightful owner should generally be able to successfully authenticate to the token. Any sort of delays resulted from the unlocking mechanism such as typos during typing or biometric false negatives may contribute to the mechanism's inability to provide this property. As an example, PINs have a limited input length

²<http://www.emc.com/domains/rsa/index.htm?id=1156>

³As discussed in the previous chapter 2, each Picosibling contains a k-out-of-n secret used to reconstruct the "Pico Master Key". The user therefore needs to choose the right combination of Picosiblings in order to unlock the Pico, which may prove difficult

and character set size which makes infrequent errors unlikely and therefore offer the property. Biometric mechanisms, based on the type and implementation may quasi-offer the property, although they generally do not.

Easy-recovery-from-loss

The meaning of this property was modified to reflect the context of token unlocking. It is offered if the user may easily recover from the loss of authentication credentials. Depending on the scheme, this may include the loss of auxiliary devices, forgotten credentials, difference in biometric features. As an example, forgotten PINs offer the property as they generally require a simple reset using an online service.

Accessible

The mechanism is usable by any user regardless of any disabilities or physical conditions. In the original paper, passwords are offered as an example of a scheme which offers this property. A gait recognition unlocking scheme would not offer this property.

Negligible-cost-per-user

The total cost per user of using the scheme, enquired by both the user and the verifier, is negligible.

Mature

A large number of users have successfully used the scheme. Any open source projects involving the mechanism, as well as any participation not involving its creators contribute to this property. For example, passwords are widely used and implemented and therefore offer the property.

Non-proprietary

Anyone can implement the token unlocking scheme without having to make any payment such as royalties. The technologies involved in the scheme are publicly known and do not rely on any sort of secret.

Resilient-to-physical-observation

An attacker would not be able to impersonate the owner of the token after observing him authenticate. Based on the number of observations required for the attacker to unlock the token, the scheme may quasi-offer the property. The original

paper suggests 10-20 times to be sufficient, although it is just an approximation. Physical observation attacks are not restricted to shoulder surfing, and may include video cameras, keystroke sounds, or thermal imaging of the PIN pad.

Resilient-to-targeted-impersonation

An attacker should not be able to impersonate the owner of the token by exploiting knowledge of personal details. This may include birthday, full name, family details, and other sensitive information. The scheme should also be resilient to pre-recordings of biometric information which may then be replayed to the authenticator.

Resilient-to-throttled-guessing

The scheme is resilient to attacks with a guessing rate restricted by the mechanism. The process cannot be automated due to the lack of physical access to authentication data. This may be achieved using tamper resistant memory. As an example, PINs offer this property because SIM cards become locked after only three unsuccessful attempts.

Resilient-to-unthrottled-guessing

The scheme is resilient to attacks with a guessing rate unrestricted by the mechanism. Even though the guessing process is only restricted by the attacker's computational power, the scheme would still not be bypassed within reasonable time. The original paper suggests that if the attacker may process 2^{40} to 2^{64} guesses per account, they would only be able to compromise less than 1% of accounts. Since tokens are generally designed to have one owner, the original description will be adapted to a single account. Therefore the property is granted only if an attacker requires more than 2^{40} attempts.

Resilient-to-theft

The property applies to schemes which use additional hardware other than the token. If the additional hardware becomes in the possession of an attacker, it is not sufficient to unlock the token. For example, auxiliary biometric devices used in the conjunction with the token offer this property. In this case the token would still not be unlocked using the hardware alone. Picosiblings however only quasi-offer the property. Although they generally rely on proximity to the Pico, the two special shares allow the owner to lock the token remotely.

Unlinkable

Using the authentication input with any verifier using the same authentication mechanism⁴ does not compromise the identity of the token owner. As an example the link between a PIN and its owner is not strong enough to make a clear link between the two. However, biometrics are a prime example of schemes which do not offer this property.

We have augmented the subset from the original UDS framework with a number of properties relevant which are relevant to Picosiblings, PINs, as well as other token unlocking mechanisms. The following list is part of the project's contributions to the overall evaluation framework.

Continuous-authentication

The token unlocking scheme re-authenticates the user periodically. The process doesn't need to be hidden from the owner, but it is required to be effortless. The token should remain unlocked and usable in the presence of its owner. The scheme needs to detect when the owner is no longer in possession of the token, and lock the device accordingly. When locked, any open authentication session managed by the security token will be closed. The concept is mentioned by Bonneau et al [5], but not included in the UDS framework. It is discussed in more detail by Stajano [7] as one of the benefits of the Pico project. Using the UDS classification of the original framework, the property belongs to the Security category.

Multi-level-unlocking

The unlocking scheme provides quantifiable feedback, not just a locked or unlocked state. The mechanism offers the possibility of supporting multiple token security permissions. These would be granted based on the confidence level that the user trying to unlock the token is its owner. For example, a 70% confidence level that the owner is present may allow the user to access an email account, but not make any sort of payments or banking transactions. Passwords only provide a "yes" or "no" answer and therefore do not offer this property. Biometric mechanisms can offer this property. Their output is either a probability or some sort of distance metric that data was collected from the owner. Different confidence levels could

⁴The authentication mechanism is not necessarily used for token unlocking. Any sort of mechanism which requires user authentication is a valid option.

therefore enable different security permissions. Using the UDS classification of the original framework, the property belongs to the Security category.

Non-disclosability

The owner may not disclose authentication details neither intentionally or unintentionally. This is a broader version of the "resilient-to-phishing" and "Resilient-to-physical-observation" properties from the original UDS framework. However, the focus here is that the token may only be used by its owner. This is an important property in enterprise situations where the security token should not be shared. Passwords and other schemes based on secrets do not offer this property as the owner could share it with another user without any difficulty. Biometric mechanisms however cannot be easily disclosed. Based on the UDS classification the property belongs to the Security category.

Availability

The owner has the ability of using the scheme regardless of external factors. The ability to authenticate should not be impaired by the authentication context such as traffic noise, different light intensities, or restricted movement space. The property is not related to physical disabilities preventing the user from using the scheme but only on contextual influences on data collection. As an example gait recognition would only function while moving on foot and therefore does not offer the property. A mechanism requiring a PIN on the other hand would work in any circumstance. Using the UDS classification of the original framework, the property belongs to the Usability category.

3.3 Example evaluation

We will demonstrate how the framework should be used by assessing three token based authentication mechanisms: Picosiblings, PIN, and Face-unlock. Each scheme represents a different type of authentication method. Picosiblings essentially are a secret the owner has, PINs are a secret the owner knows, and Face-unlock reflects who the owner is. Results in the Picosiblings section will be used in the following chapter as a benchmark for comparison with our proposed token unlocking scheme.

3.3.1 Picosiblings

Unlocking the Pico token requires k -out-of- n secrets used to reconstruct the Pico Master Key. Each Picosibling contains a secret that is transmitted to Pico using a secure connection via a radio channel. Given enough secrets the master key is reconstructed, and Pico becomes unlocked.

The scheme doesn't require from its owner any known secret and therefore is "memoryless". Since it relies on devices embedded in everyday items the "nothing-to-carry" property is satisfied. The original paper by Bonneau marks Pico as not "easy-to-learn" due to Picosiblings management, which is a characteristic of the unlocking mechanism. It is quasi-"efficient-to-use" and has quasi-"infrequent-errors" until proven otherwise. It does not offer the "easy-recovery-from-loss" property. The unlocking mechanism relies only on radio communication with the token. This makes it invariable to external factors therefore offering the "availability" property.

The original paper marks Pico as not "accessible" due to the coordinated use of camera, display, and buttons. However, the Picosiblings are "accessible" because they are embedded in everyday accessories that any user can wear. Pico doesn't aim to satisfy the "negligible-cost-per-user" property, and since no realistic Picosiblings cost estimate exists we will consider the property is not offered. The scheme is at the stage of a prototype, with no external open source contributions, and little user testing. For these reasons it is not considered to be "mature". Frank Stajano states in his paper [7] that the design of Pico and Picosiblings are not patented, and no royalties are due. The only requirement for implementing the design is to cite the paper, which makes the unlocking mechanism "non-proprietary".

Since the scheme does not rely on any user input it is "resilient-to-physical-observations". Based on the description of Picosiblings given by Stajano [7] the scheme offers the "resilient-to-targeted-impersonation", "resilience-to-throttled-guessing", and "resilient-to-unthrottled-guessing" properties. Any attacker which comes in possession of the Picosiblings may unlock the Pico. However due to the auxiliary shared secrets⁵ the scheme is quasi "resilient-to-theft". Each Picosibling only works with one verifier (its

⁵Picosiblings also relies on two special shares. One is unlocked using biometric authentication, and the other is provided by an external server. Using these shares would only grant the thief a limited time window before the token is either locked remotely or the shares expire.

master Pico), and therefore offers the "unlinkable" property. The scheme was designed to provide "continuous-authentication". Because of the k-out-of-n master key reconstruction mechanism, Picosiblings only have the locked or unlocked states and therefore do not offer "multi-level-unlocking". The scheme does not satisfy the "non-disclosability" property. The owner is free to disclose authentication credentials simply by giving their Picosiblings.

3.3.2 PIN

PINs are token authentication mechanisms similar to passwords. The difference between the two is that they use a smaller set of input characters. Additional protection comes from steep security measures when the authentication has failed. As an example, typing 3 wrong PINs on a mobile phone would lock your SIM card. A lot of the PIN properties should however be similar with those offered by passwords.

The scheme relies on knowing a secret, which is not "memorywise-effortless". It does however offer the "nothing-to-carry" property. Because of its similarity with passwords users find it "easy-to-learn". The small character set allows for fast user input and validation making PINs "efficient-to-use". Mistakes however may still occasionally occur, and due to the lack of visual feedback⁶ the scheme only quasi-offers "infrequent-errors". PINs are generally easily reset by the manufacturer using online services, granting them "easy-recovery-from-loss"⁷. The scheme offers the "availability" property, as the authentication process cannot be impaired by external factors.

Just as passwords PINs score all points in deployability. They can be used regardless of disabilities, making them "accessible". They have virtually no cost, satisfying the "negligible-cost-per-user" property". Being a subset of passwords, the mechanism is considered to be "mature" and "non-proprietary".

From a security perspective PINs score poorly. They are not "resilient-to-physical-observation". Anyone can eavesdrop the input of a PIN either by shoulder surfing or recording with a camera. Similarly to passwords, PINs are often written down in plain

⁶If existent, visual feedback for PINs generally consists of '*' characters.

⁷An example of this is the RSA SecurID. An example reset procedure is described at the following link: <http://uk.emc.com/collateral/15-min-guide/h12278-am8-help-desk-administrator-guide.pdf>

sight. However, in the lack of relevant studies⁸ we will mark the scheme to quasi-offer the "resilient-to-targeted-impersonation" property. The restricted character set makes PINs adopt harsher security policies when provided invalid input. They are generally locked after three bad attempts, making them "resilient-to-throttled-guessing". The "resilient-to-unthrottled-guessing" property is implementation dependent. However, security tokens are dedicated devices that generally have tamper resistant memory, making unthrottled guessing not possible. Any hardware PINs may require does not compromise the mechanism, therefore offering "resilient-to-theft". Users have the freedom of choosing any PIN. Even in situations when reused with multiple tokens, credentials are generally salted and therefore "unlinkable". The scheme does not offer "continuous-authentication" due to explicit requests. They can only provide locked or unlocked feedback, and therefore do not offer "multi-level-unlocking". The owner may disclose their PIN at any time, making the "non-disclosability" property unsatisfied.

3.3.3 Face unlock

Although not currently used as a security token unlocking mechanism, face recognition is a viable biometric authentication scheme. It can be ported for a token such as Pico, which is designed to have a camera. With a variety of possible implementations, for accessibility reasons we will analyse the Android face unlock mechanism.

Face unlock is "memorywise-effortless", as any other biometric scheme. It offers the "nothing-to-carry property", the camera being embedded as part of the token. The mechanism is "easy-to-learn", since it only needs the user to look at the camera. The authentication process is performed almost instantly, making the scheme "efficient-to-use". The scheme is dependent on camera positioning, obstructing objects (i.e. glasses, earrings), and face mimic. In conjunction with the UDS framework assessment of biometrics in general, the scheme does not offer "infrequent-errors". If the scheme no longer functions as a result of change in facial traits, Android has a backup unlocking mechanism. This may also be used to disable or recalibrate the scheme, therefore offering "easy-recovery-from-loss". The "availability" property is not satisfied due to the dependence on external factors such as light or obstacles.

⁸Just as Bonneau et al suggest, a relevant study would assess acquaintances' ability to guess the PIN of a subject.

Android face recognition is "accessible" for anyone regardless of disabilities. It offers the "negligible-cost-per-user" property, given that the hardware was already present in devices without face recognition features. Due to limited user exposure it is only quasi-mature". The scheme relies on proprietary software and therefore is not "non-proprietary".

Observing the owner authenticate using the scheme does not provide any advantage to an attacker. The scheme therefore offers the "resilient-to-physical-observations" property. Targeted impersonation is an issue with any biometric mechanism. The scheme is vulnerable to replay attacks (i.e. a picture of the owner's face) and does not offer the "resilient-to-targeted-impersonation" property". The "resilient-to-throttled-guessing" and "resilient-to-unthrottled-guessing" properties do not apply. Given the Android implementation, neither does "resilient-to-theft". The same authentication data is used with any verifier, and therefore the "linkable" property is not offered. The scheme is implemented without "continuous-authentication" or "multi-level-unlocking" although both can be supported by biometric mechanisms. Given the possibility of deliberately providing data for a replay attack, the scheme only quasi-offers the "non-disclosability" property.

3.4 Conclusions

We have developed a token unlocking evaluation framework. The result is strongly related to similar work by Bonneau et al [5] which was summarised at the beginning of the chapter. Some properties needed to be adapted to fit the context of a security token. We have also contributed with 4 original properties.

The framework was applied for three sample token unlocking mechanisms. A summary of the results is posted in table 3.1. Each property is highlighted with an appropriate colour in order to allow for quicker analysis. These will serve as a benchmark for the proposed solution.

As the table shows, none of the example schemes completely dominates the others. They receive mixed scores in terms of availability and security. PINs dominate in terms of deployability, receiving a perfect score.

Property	PIN	Picosiblings	Face recognition
Memorywise-effortless	Not-ordered	Ordered	Ordered
Nothing-to-carry	Ordered	Quasi-ordered	Ordered
Easy-to-learn	Ordered	Not-ordered	Ordered
Efficient-to-use	Ordered	Quasi-ordered	Ordered
Infrequent-errors	Quasi-ordered	Quasi-ordered	Not-ordered
Easy-recovery-from-loss	Ordered	Not-ordered	Ordered
Availability	Ordered	Ordered	Not-ordered
Accessible	Ordered	Ordered	Ordered
Negligible-cost-per-user	Ordered	Not-ordered	Ordered
Mature	Ordered	Not-ordered	Quasi-ordered
Non-proprietary	Ordered	Ordered	Not-ordered
Resilient-to-physical-observations	Not-ordered	Ordered	Ordered
Resilient-to-targeted-impersonation	Quasi-ordered	Ordered	Not-ordered
Resilient-to-throttled-guessing	Ordered	Ordered	Ordered
Resilient-to-unthrottled-guessing	Ordered	Ordered	Ordered
Resilient-to-theft	Ordered	Quasi-ordered	Ordered
Unlinkable	Ordered	Ordered	Not-ordered
Continuous-authentication	Not-ordered	Ordered	Not-ordered
Multi-level-unlocking	Not-ordered	Not-ordered	Not-ordered
Non-disclosability	Not-ordered	Not-ordered	Quasi-ordered

Table 3.1: Token unlocking framework sample assessment.

Chapter 4

Design

4.1 Design requirements

The framework evaluation of Picosiblings provides insight as to how the scheme can be improved. We identify as a key downside that it does not guarantee the identity of the owner. This information is mainly inferred from the number of Picosibling shares in the proximity of the Pico. However, anyone may be in possession of the shares, therefore being temporarily granted full authentication privileges. This is reflected in the evaluation by failing to fully offer the "resilient-to-theft" and "non-disclosability" properties. A further improvement can be made by introducing "multi-level-unlocking", allowing for multiple levels of authentication depending on the confidence in the owner's presence.

The Pico design proposed by Stajano [7] claims two properties that also need to be supported by the token unlocking mechanism: the authentication process is memory effortless; and the unlocking scheme needs to support continuous authentication¹.

One of the goals of designing the new Pico unlocking mechanism is to fully satisfy the the properties presented in this section.

¹Continuous authentication is defined by the ability to re-authenticate the user without the need for any physical effort.

4.2 Proposed solution

The idea explored in this dissertation project is to simultaneously use multiple continuous authentication mechanisms. Each mechanism needs to provide a quantifiable confidence level which will be used in calculating a combined score. This satisfies the memoryless and continuous authentication properties required by Pico. By combining mechanisms we achieve a higher confidence of correctly identifying the owner. Furthermore, given that each individual mechanism supports continuous authentication, using them simultaneously does not create any inconvenience for the owner.

Multi-level unlocking model

The Pico token should no longer enter a general locked or unlocked state. Its most important secret, the "Pico Master Key" should be kept in tamper resistant memory, and be accessible at all times. Using the overall score computed by the proposed mechanism, Pico should offer granular user authentication. Each app needs to be associated with a confidence level defined during the registration process. If the overall confidence of the unlocking mechanism exceeds the app's confidence level, then the token becomes "unlocked" for that specific app. All authentication sessions between Pico and apps need to be managed independently based on this model.

The scheme should achieve continuous authentication, while correctly identifying the owner of the token. For this reason we have decided that authentication mechanisms combined in the scheme need to be based either on biometrics or behavioural analysis. Biometric features that can be used with this scheme include iris, face, voice, and gait. Behavioural sources of data can be obtained from frequent GPS location, travel paths, wireless network connections, and others.

Each mechanism of the scheme is assigned a different initial weight that is based on the level of trust it offers in identifying the owner. This doesn't necessarily need to be related to the precision of the mechanism, but it would be a good indicator for choosing the value.

The solution offered in this project is different from simply stating that Pico is using biometric data as an unlocking mechanism. The novelty in the design is based on how data is combined in order to compute the overall confidence level.

Decaying weights

Data samples captured for the owner authentication process are not always meaningful. For example, accelerometer values for gait recognition are only usable when the user is travelling on foot. Depending on how the sensors are integrated with the Pico, camera input for face recognition may not always capture a valid image. The confidence of each mechanism should therefore decrease in time from the last valid authentication sample. This introduces another original feature of this scheme, which is having a decaying weight. Each mechanism starts with a predefined initial value, reflecting the weight of the mechanism in the overall unlocking process. This value decreases in time until a valid user data sample is provided to the mechanism for authentication.

Let us take for example a voice recognition mechanism which samples data every minute. The current weight of the mechanism is 0 so its output is completely ignored. The next sample is recorded, and the voice recognition mechanism outputs a confidence of 70% that the owner is present. After the successful recording, the mechanism weight is updated to its predefined starting value of 30. For the next 10 minutes the owner will be silently reading a book. Since the mechanism only identifies background noise, the weight value of 30 decreases in time. This will induce a smaller impact of the voice recognition mechanism on the overall score. Each individual confidence level of a mechanism can decrease up to 0, at which point the mechanism is ignored. Computing the overall score will be explained in more detail later in the chapter.

4.2.1 Explicit authentication

We need to consider the case where the owner wants to use Pico to authenticate to a high security app, given a low confidence level from the authenticator. As an example, the Pico owner wants to access their bank account after sitting silent in a dark room for the past hour. Let us say the app requires a confidence level of 95%. Due to the lack of valid authentication data, the authenticator only outputs a 20% overall

confidence that the owner is present. To solve this problem we have introduced the concept of explicit authentication mechanisms. When the confidence score drops below the threshold required by an app, the user is given the chance to authenticate through an explicit request to provide valid sample data to one or more mechanisms.

Combining explicit and continuous authentication can be performed consistently with the current design. Whenever explicit authentication is required, the only difference for the owner is that they become aware of the authentication process. This grants them the chance to produce valid input for the authenticator, which may then proceed to compute an accurate score. Given that prior to the explicit authentication request the unlocking mechanism didn't produce a high enough confidence, it is assumed that this will also happen prior to that. Therefore, explicit authentication requests need to have a slower decay rate. This will enable the continuous authentication process.

Authentication feedback

Each mechanism outputs a value, which is the probability that the sample data belongs to the owner of the token. Upon each recording, this probability is updated using Bayes' Law. This process is also known as a Bayesian update, and is described in the following equation:

$$P(H|E) = \frac{P(H) * P(E|H)}{P(E)} \quad (4.1)$$

In the equation above:

- E: Stands for evidence and in this case represents the data sample.
- H: Stands for hypothesis. In this case we refer to the hypothesis that the owner is present.
- $P(H|E)$: Represents the probability of hypothesis H after observing evidence E . This is the final probability we are trying to compute after each sample. It is also known as the posterior probability.

- $P(H)$: Represents the probability of hypothesis H before observing evidence E . This is also known as the prior probability and is the probability computed at the previous step.
- $P(E|H)$: Represents the probability that the current evidence belongs to hypothesis H . It is the probability outputted by the mechanism given the sample data.
- $P(E)$: This is the model evidence, and has a constant value for all hypothesis.

Although $P(E)$ is constant we need its value in order to calculate $P(H|E)$. We can compute it using the "Law of total probability", which is the following:

$$P(E) = \sum_n P(E|H_n) * P(H_n) \quad (4.2)$$

Using equation 4.2 Bayes' Law equation 4.1 becomes:

$$P(H|E) = \frac{P(H) * P(E|H)}{\sum_n P(E|H_n) * P(H_n)} \quad (4.3)$$

Our model however, contains only two hypothesis²: the recording of the data either belongs to the owner, or not. We can therefore consider $P(H)$ to be the probability that the data belongs to the owner and $P(\neg H)$ the probability that the data belongs to someone else. This means the value of $P(\neg H)$ is $1 - P(H)$ and $P(E|\neg H) = 1 - P(E|H)$. Introducing this in equation 4.3, the rule for updating the mechanism's probability becomes:

$$P(H|E) = \frac{P(H) * P(E|H)}{P(H) * P(E|H) + P(\neg H) * P(E|\neg H)} \quad (4.4)$$

Equation 4.4 represents the final probability that the owner is present given the sampled data. All the variables in this equation are known, for reasons explained above.

Thus far we have defined how individual scores are calculated, and that each mechanism has a decaying weight. Using this data we can calculate the overall score of the scheme.

²Arguably there is a third case where the data sample is not a valid recording of an user. This case is ignored and no probability is computed. The only result in this scenario is the resuming of the decay process in the weight of the mechanism.

This is performed quite trivially by using a weighted sum. The following equation describes the process.

$$P_{Total} = \frac{\sum_{i=1}^n (w_i * P_i(H|E_i))}{\sum_{i=1}^n w_i} \quad (4.5)$$

Pico needs to use the overall score to adjust the state of its authentication sessions. If the score required by an app is lower than the overall score provided by the scheme, the user is granted authentication access. To provide continuous authentication, Pico will make periodic requests for confidence level updates. Feedback is provided by the unlocking mechanism based on available sample data, and decaying weights.

4.3 Related work

Clarke et al [13] present statistics confirming the need for an unlocking scheme different from PINs. They conduct a couple of surveys trying to assess the reliability of a PIN in unlocking a mobile phone. The study involves 297 participants and assess: day to day usage of mobile phone devices, existing authentication mechanisms, and the users' attitude towards further security options. The paper reveals a high number of bad practices involved in PIN authentication: 45% of owners never change the default factory code, 42% only change it once after buying the device, reusing the PIN with other authenticators, forgetting the PIN, and sharing the PIN with someone else.

A promising result showed in the paper is that 83% of users are willing to accept some sort of biometric authentication mechanism in order to unlock their devices. The following biometric mechanisms were included in the study: fingerprint analysis, voice recognition, iris recognition, hand recognition, keystroke analysis [14], and face recognition. The paper also shows that 61% of users would accept a non intrusive biometric continuous authentication mechanism. Using multiple biometrics for continuous authentication is mentioned briefly, but each mechanism is used individually based on what the user is doing. As an example, when the user walks he is authenticated using gait recognition, and while he is speaking on the phone, voice recognition. This is a divergence point from what we are trying to achieve in this dissertation.

In a different paper, Clarke et al [6] study PIN alternatives for mobile phone unlocking. The authors conduct a survey with interesting results. A remarkable 11% of participants were not even aware of PIN authentication. An average of 81% of participants agree that PINs should be replaced with a mechanism that provides better security. Although they report the need and desire for a different type of phone unlocking, many of them do not use currently available alternatives.

Gregory Williamson [15] writes in his PhD dissertation about the need for an enhanced security authentication mechanism for on-line banking. He proposes a multi-factor authentication model, and presents two interesting options: the traditional one where both mechanisms are required in the multi-factor model (blanket authentication), and one where the second authentication mechanism is only requested from the user if the transaction appears to be risky (risk mode authentication). A risky situation is defined as either an important transaction like withdrawing money, or a transaction made under unusual circumstances such as using an unknown device.

A similar approach to "risk mode authentication" presented by Williamson [15] is proposed in this project. Our scheme yields a confidence level which may or may not be sufficient to allow access for different transactions. If the confidence level is not high enough, an explicit authentication mechanism will prompt the user for input. As the dissertation by Williamson shows, 75% of users questioned in his study agree with having biometric authentication as a second factor authentication to passwords. This shows promising results in adopting our scheme for token unlocking purposes.

Elena Vildjiounaite et al describe in their paper [16] a similar mechanism of combining biometric authentication data on mobile phone devices. The authors identify the security risk of granting long-term authentication after a single verification challenge. They explore an alternative based on a two stage "risk mode authentication" [15]. The first stage combines biometric data in order to achieve continuous authentication. This is achieved by training a cascade classifier to a target false acceptance rate (FAR)³. Data from mechanisms is merged using a weighted sum fusion rule. Mechanism weights are chosen based on their error rates. The second stage is only enabled if the cascade classifier does not identify the owner as being present. In low noise scenarios, continuous authentication is achieved without the need for an explicit challenge 80% of the time. In

³The false acceptance rate is the equivalent of false positive precision. It is the probability of incorrectly granting authentication privileges to an user

noisy situations (city and car noise), the percentage drops ranging from 40 to 60%. The cascade classifier was trained with a FAR of 1%, with results showing a false rejection rate (FRR)⁴ of only 3 to 7%.

The paper by Elena Vildjiounaite et al [16] is similar with the solution proposed in this dissertation. It also combines multiple authentication mechanisms, each being assigned different weights. Differences between the two are in the fact that weights are maintained static in time. The weights of the sums are computed differently, and there is no mention of Bayesian updates or probabilities. Furthermore, the authors use a classifier instead of producing a confidence level, which cannot be used for granting different levels of security. The results presented by this paper are however encouraging, showing that continuous authentication presents good results using multiple biometric authentication mechanisms.

4.4 Conclusions

We have defined a conceptual model for a new Pico unlocking mechanism. The scheme is guaranteed to improve on the existing Picosiblings solution at least by offering a better way of correctly identifying its owner. We have offered reasonable explanation that it supports Pico's claims for continuous and memory effortless authentication.

An evaluation of the scheme is not yet offered, because mechanisms such as "Negligible-cost-per-user" are implementation dependent. The next chapter will present a prototype solution that will be evaluated using the token unlocking assessment framework. The results will be compared with the current Picosiblings implementation allowing for further analysis and conclusions.

⁴The false rejection rate is the probability of incorrectly denying access to the rightful owner.

Chapter 5

Implementation Prototype

In this chapter we will develop a prototype for the scheme proposed in section 4.2. Just as presented by Stajano [7], smart phones offer a cheap alternative to prototype and test. Therefore, we have chosen as an implementation platform the Android Nexus 5 handheld device. It offers enough sensors to perform biometric and behavioural analysis¹. These resources will be used to demonstrate that the scheme can be implemented using similar dedicated hardware that may offer more security features.

5.0.1 Android development and security

To gain a better understanding of different design decisions and limitations of our implementation, we will present a brief literature review of the Android development platform. Mechanisms and components will be described with an emphasis on security. The prototype developed for this dissertation is a proof of concept. However, we still aim to make proper use of available security mechanisms.

William Enck et al [17] offer a good introduction to Android application development. They focus on the security aspects of the platform. It is a relatively old paper (2008), from the same year of the Android initial release. However, the fundamental design principles and security concepts that are discussed did not change significantly. The platform's open standards were made public in November 2007. This allowed researchers such as the authors of this paper to perform a pre-release analysis of the system.

¹The full range of sensors supported by the Android platform can be found here: http://developer.android.com/guide/topics/sensors/sensors_overview.html (accessed on 28.05.2014)

Android uses as a core operating system a port of the Linux kernel. This introduces to the platform some of the Linux security mechanisms (i.e. file permissions, access control policies). On top of the kernel there is an application middleware layer composed out of the Java Dalvik virtual machine, core Java application libraries, as well as libraries which offer support for storage, sensors, display, and other device features. Applications are supported by the middleware and developed using the Android Java SDK.

The Android development model is based on building an application from multiple components. Based on their purpose, the SDK defines four types of components: activity, service, content provider, and broadcast receiver. For the purpose of brevity we will not discuss each individual component². To allow meaningful interaction, Inter Component Communication (ICC) is enabled using special objects called Intents.

The application we are developing needs to perform most of its processing in the background. It does not require any explicit user interaction. According to the Android model, this should be achieved using Services. To enable convenient component interaction, services may become bound for message passing. An important note made in the paper is that while a Service is bound, it cannot be terminated by an explicit stop action. This provides an useful guarantee regarding its lifetime that we will use in the prototype.

The paper discusses two types of Android security enforcements: ICC and system level. System level security is based on the Linux permission model. Each app is associated with an UID and GID. This allows internal storage access control restrictions, keeping application data sandboxed from other applications.

ICC security is the main focus of the paper. Intent communication is based on commands sent to the `"/dev/binder"` device node. The node needs to be world readable and writeable by any application. Therefore, Android cannot mediate ICC using the Linux permissions model. Security relies on a Mandatory Access Control (MAC) framework enforced by a reference monitor. This mechanism validates requests sent to the `"/dev/binder"` node.

During development, each application needs to define a manifest file³. Some of the security configurations defined in this file are: declared components and their capabilities, permissions required the app, and permissions other apps need to have in order to interact with app components. These entries are used as labels for the MAC framework.

²More details on the role of each component can be found on the Android website: <http://developer.android.com/guide/components/fundamentals.html>

³Full details regarding the manifest file can be found on the Android website: <http://developer.android.com/guide/topics/manifest/manifest-intro.html>

Using the app manifest file, each component can be defined as either public or private. This refinement is configured by the `\exported` field. It defines whether or not another application may launch or interact with one of its components. When this paper was written, the `\exported` field was defaulted to `\true`. However, as shown by Steen and Mathias [18] in 2013, starting with Android 4.2 the default of this value was changed to `\false`, and now conforms to the `\principle of least privilege`.

Components listening for Intents need to have an intent-filter registered in the application manifest file. This allows them to export only a limited set of intents to other applications. Further restrictions to Intent objects are offered by the SDK using permission labels. This mechanism provides runtime security checks for the application. It is an additional prevention mechanism for data leaks through ICC. An application may broadcast an event throughout the system. By using permission labels, only applications that have the respective permission may have access to the event. Furthermore, Services may check for permissions when they are bound by another component. This allows them to expose different APIs depending on the binder.

Steen and Mathias [18] focus on deeper issues of the Android platform. They show and how problems are solved from one Android version to the other. However, OEMs tend not to update the software of their devices once they have shipped, which poses a number of security issues.

The starting point of understanding Android security and how it is bootstrapped is the five step booting process:

1. Initial bootloader (IBL) is loaded from ROM.
2. IBL checks the signature of the bootloader (BL) and loads it into RAM.
3. BL checks the signature of the linux kernel (LK) and loads it into RAM.
4. LK initialises all existing hardware and starts the linux `\init` process.
5. The init process reads a configuration file and boots the rest of LA.

The android security model is split by the paper in two categories: system security, and application security.

Android provides a keychain API used for storing sensitive material such as certificates and other credentials. These are encrypted using a master key, which is stored using AES encryption. Security needs to begin somewhere. An assumption has to be made about a state being secure from which multiple security extensions can be made. In this case, the master key is considered to be that point of security. However, given a rooted

device, the master key itself may be retrieved from the system therefore compromising all other credentials. The Android base system (libraries, app framework, and app runtime) is located in the `\system` partition. Although it is writeable only by the root user, exploits which grant this privilege exist.

From the user's perspective, an interesting "feature" which may affect the flow of information within Android is the fact that applications from the same author may share private resources. When installing an application the user needs to accept its predefined set of permissions. Due to resource sharing, a situation may present itself when an application that has permissions for the owner's contacts may communicate with an application that has permissions for internet in order to leak confidential data. A developer may therefore construct pairs of legitimate applications in order to mask a data flow attack.

The Android OS offers a number of memory corruption mitigations in order to avoid buffer overflow attacks, or return oriented programming. The following list presents these low level security mechanisms:

- Implements `mmap_min_addr` which restricts `mmap` memory mapping calls. This prevents NULL pointer related attacks.
- Implements XN (execute never) bit to mark memory as non-executable. The mechanism prevents attackers from executing remote code.
- Address space layout randomisation(ASLR) implemented starting with Android 4.0. This is a first step to preventing return oriented programming attacks. The memory location of the binary library itself is however static. After a number of attempts using trial and error, the attacker may succeed using return oriented programming.
- Position independent and randomised linker(PIE) is implemented starting with Android 4.1 to support ASLP. This makes the memory location of binary libraries to be randomised.
- Read only relocation and immediate binding space(RELro) was implemented starting with Android 4.1. It solves an ASLR issue where an attacker could modify the global offset table (GOT) used when resolving a function from a dynamically linked library. Before this update an attacker may insert his own code to be executed using the GOT table.

A number of application security mechanisms are in place to make Android a safer environment for its users. A program also known as a "on device Bouncer" prevents malware

to be distributed from within the Android App store (Google Play). The purpose of the bouncer is to verify apps prior to installation by checking for malware signatures and patterns. Secure USB debugging was introduced starting with Android 4.4.2. This only allows hosts registered with the device to have USB debugging permissions. The mechanism is circumvented if the user does not have a screen lock.

The paper presents that the Android OS is responsible for 96% of mobile phone malware. The authors claim that this is the case due to 4 big issues of the Android platform:

1. Security updates are delayed or never deployed. This is due to a number of approvals that an update needs to receive. This introduces an additional cost to the manufacturer (OEM), that does not generate any revenue. The majority of teams working on the Android platform are focusing on current releases. In most cases there are simply not enough resources to merge Google security updates to the OEM repository. Furthermore, the consequences of a failed OS update may cause "bricking" of the device, which is a huge risk for the manufacturer. All these issues lead to very few security updates. Therefore, important features such as RELro are never deployed, making older Android releases vulnerable.
2. OEMs weaken the security architecture and configuration of Android by introducing custom modifications before they roll out a device.
3. The Android permission model is defective. According to Kelley et al [19], most users do not understand the permission dialogue when installing an application. Furthermore, even if they could understand the dialogue, most of the time it is ignored in order to use the new exciting app. According to the same study, most applications are over-privileged. This is due to developers not understanding what each privilege grants. Furthermore, as previously pointed out, applications from the same owner may share resources and implicitly privileges. This creates a valid information flow attack path.
4. Google Play has a low barrier for malware. A developer distribution agreement (DDA) and a developer program policy (DPP) need to be agreed to and signed by the developer before submitting the application to the Android market. However, Google Play does not check upfront if an application adheres to DDA and DPP. The application is only reviewed if it becomes suspect of breaking the agreements. Furthermore, according to [20] there are ways of circumventing the Bouncer program⁴.

⁴An example of such an application is presented in an article written in Tech Republic: <http://www.techrepublic.com/blog/google-in-the-enterprise/malware-in-the-google-play-store-enemy-inside-the-gates/#> (visited on 29.05.2014).

We have briefly presented the Android development model, existing mechanisms, and the security of the platform. Given this information we may proceed to present the prototype implementation of the token unlocking scheme.

5.1 Authenticator design

The Android user authenticator prototype is designed to work as a bound service implemented in the `UAService`⁵ class. The service collects data from each mechanism and computes the final authentication confidence level. The result is provided to clients either after an explicit request or through periodic broadcasts.

Each authentication mechanism may have a different requirement for sampling and processing data. As an example, voice recognition may gather optimal data during a phone call⁶, while face recognition when the phone screen is unlocked. Therefore, to enable more flexibility in the individual mechanisms' implementation, we have chosen each to be developed as an independent service.

`UAService` communicates with the authentication mechanisms by binding their service. This allows message passing using a `ServiceConnection` object. On predefined time intervals `UAService` acquires the confidence level and weight of each mechanism. Using this data it then calculates the overall result according to the design in section 4.2.1. Feedback is sent back to each registered client for interpretation.

5.2 Implementation details

This section presents the implementation of the scheme proposed in section 4.2. The full implementation of the prototype can be downloaded from from github:

`\https://github.com/cristiantoader/fyp-pico".`

5.2.1 Main application and services

The Android token unlocking scheme is implemented as a bound service in the `UAService` class. According to the Android documentation a bound service is the server in the client-server interface. It enables other components to send requests and receive responses by

⁵The name of the class stands for User Authentication Service

⁶Call events can be intercepted by registering a listener for the `PHONE_STATE` event

binding to it. It is developed as a regular service that implements the `\onBind()` callback method to return an `IBinder`.

According to the Android development API guide⁷ there are two independent scenarios describing the lifetime of a bound service, depending on how the service was started:

1. If the service was not previously running, and a `\bindService()` command is issued by a component, the service is kept alive for as long as clients are still bound. A client becomes unbound by calling `\unbindService()`.
2. If the service is started using `\onStartCommand()` it can only be stopped if it has no bound clients and an explicit request is made either via `\stopSelf()` or `\stopService()`. Unlike the previous case, its lifetime persists even with no bound clients.

The prototype we have developed takes into account the lifetime of bound services. The `\PicoMainActivity` class calls `\startService()` and `\bindService()` to the `\UAService` component. When `\PicoMainActivity` gets sent to background and loses control of the screen, the service is not explicitly unbound. This should prevent other components from killing the service.

A safer alternative would be to create a root service. This requires modifications to the system partition. The process does not resume to simply gaining root privileges and making the modifications. The root directory is mounted as ramdisk, and therefore any direct changes will be reverted once the device is rebooted. In order to make persistent modifications, the user needs to alter the boot image and re-flash it on the device. The service needs to be compiled using the Android NDK C compiler. The binary has to be included in the system partition of the boot image in order to be run at start-up by the `\init` process. The `\init.rc` configuration file used by `\init` also needs to be configured to start the service.

`\UAService` is a central node in the application. It gathers data from the authentication mechanisms, computes the overall confidence data, and sends feedback back to the clients.

In order to receive authentication updates, clients need to bind `\UAService`. Communication is enabled using the `\Messenger` interface, which is the simplest way to perform Inter Process Communication (IPC). The `\Messenger` queues all requests on a single thread, and therefore the application does not require to be thread safe. The following commands are exposed to clients such as Pico through the `\what` parameter of the `\Message` class:

⁷<http://developer.android.com/guide/components/bound-services.html>

MSG_REGISTER_CLIENT

Used for registering a client for periodic broadcasts⁸ of the current unlocking confidence level. Feedback is provided at a fixed time interval of 1000ms.

MSG_UNREGISTER_CLIENT

Used for any application to unregister as a listener of this service. No additional parameters required.

MSG_GET_STATUS

Used by any application when an authentication request is needed. Although the service periodically broadcasts to its registered clients what is the authentication status, explicit requests may also be performed using this \Message".

The communication between \UAService" and an authentication mechanism service is intermediated by an \AuthMech" object. Each \AuthMech" is responsible for interfacing the communication with its corresponding service. From the authentication mechanism's side, each service needs to extend the \AuthMechService" abstract class. This standardises the communication with the \AuthMech" objects, and implicitly \UAService". This software engineering approach facilitates adding additional mechanisms with minimal changes to the original code.

Each \AuthMechService" is implemented as a bound service. When binded by \UAService" through \AuthMech" they expose the following message passing interface:

AUTH_MECH_REGISTER

Used for registering the \UAService" client to the \AuthMechService".

AUTH_MECH_UNREGISTER

Used for unregistering the \UAService" client to the \AuthMechService".

AUTH_MECH_GET_STATUS

Used by the \UAService" to request the authentication feedback from the \AuthMechService". The Message response has in \arg1" the authentication confidence multiplied by the decayed weight, and in \arg2" the original weight of the mechanism.

⁸An alternative implementation explored in the project was to have each client also register a confidence level using the "arg1" parameter of a "Message". In this case, the authenticator would only provide each client with a locked/unlocked result. However, this would shift the meaning of client to that of an authentication session, with state managed by the unlocking scheme. A client would therefore have multiple connections, requiring more IPC. Since all "Messenger" requests made to "UAService" are queued to a single thread, this would slow down the feedback process and possibly even lead to a denial of service attack. Therefore we have chosen to reduce the communication overhead, and have each client manage the status of its authentication sessions based on the confidence level provided by the unlocking scheme.

5.2.2 Authentication mechanisms

In order to create a functional prototype of the scheme, we have implemented a number of user authentication mechanisms. We will not focus on the quality of the biometric and behavioural mechanisms. Their sole purpose is to demonstrate that the design of the scheme is functional, and can be implemented using only a smart phone.

A wide range of biometric mechanisms can be developed using sensor data offered by Android devices. Examples include face recognition, voice recognition, iris scanning, keystroke analysis, gait recognition, and many others. The requirements when developing an authentication mechanism for the scheme are:

1. The result needs to be quantifiable, preferably as a percentage ranging from 0 to 100.
2. The mechanism needs to support continuous authentication of the user.
3. The authentication process needs to be effortless and preferably unobtrusive for the user.

Android provides an extensive sensor API that can support the scheme. This can be used to develop a number of continuous authentication mechanisms. We have listed the following non-exhaustive set of examples:

Face recognition

The mechanism is based on capturing an image of the user's face and performing face recognition. Sampling valid face images can be performed without explicit requests by predicting user behaviour. We will use as an example an user that owns a phone with a front-facing camera. When the owner is unlocking the phone, there is a high probability that they will be looking towards the screen. This provides a good opportunity for the face recognition service to capture a valid sample. Using the Android API, this can be achieved by registering a "BroadcastReceiver" to listen for the one of the following events: ACTION_SCREEN_ON, ACTION_SCREEN_OFF, or ACTION_USER_PRESENT. Based on this data, the mechanism may continue by performing face recognition based on collected data and a previously recorded sample of the owner. A simple face recognition mechanism was also implemented as part of the prototype.

Voice recognition

A voice recognition mechanism may record data either periodically, or based on Android events. It may then perform voice recognition and provide a confidence

level of the owner being present. Voice sampling does not necessarily imply a voice password. An analysis can be performed using feature extraction. This facilitates the sampling process, which may be performed at any time. With a frequent sampling period, the owner of the device is likely to be recorded while speaking, which would provide a valid data sample. For even better confidence the mechanism can be implemented to start recording when a call is either made or received. On Android this can be achieved by listening for a `PHONE_STATE` event. A simple voice recognition mechanism was implemented as part of the prototype.

Iris scanning

Similar to face recognition, this can be implemented by taking advantage of user behaviour while using the phone. When the phone is unlocked, the user is very likely to face the front camera, allowing for a good capture. The only problem with this mechanism is the quality of pictures offered by most phones. If the sampling quality is not sufficiently good, meaningful features from the iris may not be extracted. This would make the confidence level of the mechanism relatively low.

Keystroke analysis

This mechanism was inspired from a paper by Clarke et al [21]. The principle of keystroke analysis is based on the patterns in which the user types on his mobile phone. Different features can be extracted here, such as: letter sequence timings, words per minute, letters per minute, frequent used words, and others. Using this data a confidence level can be generated. This is harder to implement using solely the Android SDK. A good starting point would be to have a keyboard app developed for the user that also communicates with the authentication mechanism. If the keyboard is disabled by an attacker this should be considered, especially if the authenticator was originally configured to listen for input.

Gait recognition

This mechanism is based on analysing individual walking patterns. According to data presented by Derawi et al [22], error rates⁹ may vary between 5% to 20%. In the lack of a reliable library, efforts to implement this mechanism for Android were unsuccessful. Android offers native recognition support for walking, driving, or standing still. Applications can register a sensor callback for the `TYPE_STEP_DETECTOR` composite sensor. Whenever such an event is detected, data can be recorded from the accelerometer and validated using an algorithm similar to the one described by Derawi et al [22].

⁹ .

Ear shape analysis

Research shows (i.e. Burge et al [23], Mu et al [24]) that the shape of the human ear contains enough unique features to perform biometric authentication. Taking advantage of user behaviour, valid data can be captured and analysed using a smart phone. We suggest that a picture is taken a few seconds after a phone call event is detected. If no peripherals are attached, the user is likely to move the device towards the ear. Images captured by such a mechanism could then be used to calculate an accurate confidence level of the user's identity. This method was not tested, so therefore we cannot ensure whether the auto-focus of the camera is sufficiently fast to obtain a valid image.

Proximity devices

This is an original idea based on providing a confidence level depending on the presence of known devices. The mechanism should connect with other devices that are also running the authenticator. The two owners don't necessarily need to know one another for the acknowledgement to be performed. Whether regular travel schedules, or working in an office, users are constantly being in the presence of other known devices. This should provide a confidence as to whether the device is in the presence of its owner.

The authentication works by seeking connections with other devices. Whenever a device is identified, its ID is recorded. The mechanism needs to keep track of the number of times it has connected with another device. Some connections may be established for the first time, and should not bring any confidence. Other connections, such as the Pico of a co-worker, would probably have a high number of connections, and therefore the mechanism should output a higher confidence level in its presence. This mechanism is similar to the Picosiblings solution, but with no k-out-of-n secrets. Each Pico is essentially a Picosibling for another Pico, with each device having a difference weight based on familiarity.

As an example, when travelling with your family on holiday most of the devices there are unknown. However, given that a number of frequent IDs are in the proximity of the authenticator, the mechanism should still consider to some extent that it is in the possession of its owner.

The mechanism can be circumvented in the scenario where co-workers or friends try to unlock the Pico. Due to this downside, it should never have sufficient weight to unlock the token on its own. However, in combination with other mechanisms it would provide a good approximation of whether it is in the possession of its owner. If the device is "in good company" there is a good chance the owner is also present.

Location data

This mechanism is similar to "Proximity devices" and much easier to implement. Based on Android GPS and network location data, the phone may detect whether it is in an usual location or not. Just as "Proximity devices" this should not carry a high weight in the scheme, especially since it would not provide accurate results in scenarios such as holidays.

Service utilisation

This mechanism exploits patterns in the Android phone's service and app utilisation. Based on current running applications, services, and the time they were started we may create a model where some confidence is given regarding the ownership of the device. This mechanism would only be effective in detecting sudden changes. It would have a low weight in the overall scheme due to its lack in precision.

Picosiblings The original Picosiblings mechanism may also be used with this scheme.

Although not part of the standard set of Android device sensors, if available, a Picosiblings implementation may be included as one of the authentication mechanisms.

A number of continuous authentication mechanisms may also be used for explicit authentication. The user can be notified to provide accurate information for the following mechanism: face recognition, voice recognition, iris scanning, keystroke analysis, gait recognition, and ear shape analysis. This increases the opportunity for a valid data sample to be collected.

A number of explicit authentication mechanisms which do not satisfy the continuous authentication property of Pico may be implemented for the Android platform. It is important to note that additional mechanisms not included in this list need to satisfy the memorywise-efficient property of the token unlocking framework [3.2](#). We suggest the following mechanisms for implementation:

Fingerprint scanner

Devices that incorporate a fingerprint scanner (such as the iPhone 5S) can use the sensor as an explicit authentication mechanism. It cannot be used for continuous authentication because the user doesn't come in contact with the sensors on a regular basis. A mechanism can therefore request explicit fingerprint data, which would then be compared with the owner's biometric model, outputting a confidence for the authentication. The result will be combined in the overall scheme just as any other mechanism. The only difference will be in terms of weight and decay rate.

Hand writing recognition

The user may be prompted to use the touch screen in order to write a word of his choice. This would guarantee the memorywise-errorless property because the user doesn't need to remember any secret. The handwriting would be analysed with a predefined set of handwriting samples in order to compute the confidence level that the owner produced the input.

Lip movement analysis

According to Faraj and Bigun [25], analysing lip movement while speaking can be used for authentication. The user would be prompted to provide a data sample such as reading a word provided by the authenticator. Using lip movement authentication, a quantifiable confidence level would be produced. This mechanism can also be implemented as a continuous authentication mechanism. However, data sampling would likely have a low success rate as users tend not to have their mouth within the range of the camera.

In order to have a functional prototype of the scheme, we have developed a number of authentication mechanisms. The following sections will prove that the Android platform offers sufficient functionality for supporting the scheme. Furthermore, this should stand as proof that the design can be implemented on a dedicated Pico device with a similar set of sensors.

The following mechanisms have been implemented as part of the prototype: voice recognition, face recognition, location analysis, and a dummy mechanism used for testing. The following sections will provide details regarding their functionality and implementation process.

5.2.2.1 Dummy mechanism

A dummy authentication mechanism was developed for testing the overall scheme. It produces random confidence levels within a predefined range, which provides a good controlled environment for testing purposes.

The mechanism was developed consistently with the application model. It is implemented in the `\AuthDummyService` class, extending the `\AuthMechService` abstract class. This makes it an independent bound service with a predefined communication interface.

All authentication mechanism services have a data access object (DAO) responsible for interfacing with imported libraries and managing authentication data. In this case the

DAO only produces random confidence levels within a given range. A thread belonging to the service makes periodic requests to the DAO. This mimics an authentication mechanism that periodically samples for data. The service is updated based on the produced value.

When `\UAService` needs to update its overall confidence, it makes an `AUTH_MECH_GET_STATUS` request to the service. The reply contains the most recent confidence level multiplied by the current decay factor, and the original weight of the mechanism.

5.2.2.2 Voice recognition

The voice recognition mechanism is implemented in the `\VoiceService` class. It extends the `\AuthMechService` abstract class that defines its communication interface. When the service is created, the `\onCreate()` method is called automatically by the Android platform. The method was developed to start a thread that periodically gathers data from the microphone, performs biometric authentication, and produces a confidence level.

The library used for voice recognition is called `Recognito`¹⁰ and was developed by Amaury Crickx. It is a text independent speaker recognition library developed in Java. We do not claim that it is the best voice recognition library. However, it was best suited for the purpose of this prototype. Porting the library for Android required minimal changes. It claims very good results in scenarios with minimal background noise¹¹.

In order for the application to compile the `Recognito` library, a subset of the Java (SE) `rt.jar` was required. This is due to `\javax.sound.*` packages included in the library that are not available on Android. Trying to import and use the `\javax.sound.*` package is not possible due to the name collision with the `\javax.*` system library available on Android. Therefore, we had to include `\rt.jar` as part of the application, but without actually using it. This was purely done to trick the Android Java compiler to package the application. Using the library would generate a runtime error. This was avoided by only using raw features which require direct sound input, without any knowledge of sound file formats.

In order to gather and manage samples compatible with the `Recognito` library we have created the `\VoiceRecord` class. This is responsible of gathering microphone input using the following predefined configuration:

¹⁰The library can be downloaded using github from the following link: <https://github.com/amaurycrickx/recognito>

¹¹It was tested by the author on TED talks, such as: <https://www.ted.com/talks/browse> (visited on 06.01.2014)

- Sample rate: 44100
- Channel configuration: `AudioFormat.CHANNEL_IN_MONO`
- Audio format: `AudioFormat.ENCODING_PCM_16BIT`

The minimum buffer size required by the `\VoiceRecord` class is device dependant and pre-calculated in the constructor. The class wraps an `\AudioRecord` object used for gathering microphone data. Due to limitations of the SDK, the recording is saved as a file and loaded into memory when needed. This makes the implementation not efficient, but it does serve the purpose of the prototype.

A DAO class is created to facilitate the interface to the Recognito library. When initialised, it loads the owner configuration, and a predefined set of background noises. It then creates a Recognito object and trains it using the data. This is done using the library's `\createVocalPrint` public method.

Every predefined time interval, the `\VoiceService` authentication thread records data in `\double[]` format using the `\VoiceRecord` class. It then calls the `\recognize` public method of the Recognito object. This returns the Euclidean distance to the closest match, which is either the owner, or one of the background noises used for training.

To convert the Euclidean distance to a percentage confidence level, we define an acceptable threshold. Any result above the threshold is considered too high and is truncated to the value of the threshold. Using the following formula we convert the Euclidean distance value to a confidence level. The final result is $P(E|H)$ (the probability that the evidence belongs to the hypothesis) used in equation 4.4.

$$P(E|H) = 1 - \frac{\text{distance}}{\text{THRESHOLD}}$$

Dividing the distance over the threshold yields a confidence value between 0 and 1, where 1 is a very large distance and hence a bad result. By using one minus this value we invert the meaning. Values will range between 0 and 1, and 1 corresponds to a confidence level of 100%.

Having known $P(E|H)$ we continue to calculate $P(H|E)$ by using the Bayesian update formula defined in equation 4.4. When calculating the final confidence level of the mechanism, we multiply $P(H|E)$ with the current decay rate. Due to the message passing API, this value needs to be an integer and is therefore multiplied by 100. The overall result is stored in the service and updated whenever the decaying weight is modified. When a request is made by `UAService`, the value is returned using the `IBinder` message passing mechanism.

5.2.2.3 Face recognition

The face recognition mechanism was implemented in the `\FaceService` class, which extends the `\AuthMechService` abstract class. When created, the service starts a thread that periodically collects data from the camera. Each sample is analysed using a face recognition library, and a confidence level is outputted for the current sample. The result is multiplied by the mechanism's weight which is a decaying factor.

The library used for face recognition is a port of the `\Javafaces` library [?]. This was the closest functional library found that was compatible with the Android API. `Javafaces` is written entirely using Java (SE). Unfortunately, it makes use of the `\javax.imageio` package which is not available in the standard Android SDK. Therefore, a considerable amount of code needed to be ported for the Android platform. The new library is publicly available at the following link: [?]. It is currently not optimised for public use.

We will briefly present the changes made when porting the `\Javafaces` library. The `\BufferedImage` class had to be replaced by its Android closest equivalent, which is `\Bitmap`. All `\BufferedImage` references in the original project had to be adapted. Furthermore, The API was modified to support direct `\Bitmap` input in order to add more flexibility and lighten the main code of the authenticator.

In the original `\Javafaces` library, data formats for black and white images were assumed to have a single colour channel representing the grey value. This had to be changed to reflect the `Bitmap` convention, where all 3 colour channels are present but have the same value. Additional modifications were required due to data type mismatches, as well as other related issues.

Unfortunately, the library combined with the Android SDK does not provide accurate results. The reason is that the library requires as input a rectangle bitmap perfectly containing the face of an individual. Unfortunately, although the Android SDK offers face detection, it only provides the location of the midway coordinate between the eyes, and the distance between the eyes. Using this data alone, an accurate crop cannot be made. As a solution, yet another library would need to be used in order to properly detect face regions. This would provide better input data and would increase the precision of the mechanism.

Every predefined time interval, the authentication thread running within the `\FaceService` object samples data from the camera. This is performed by using a `\CameraUtil` object. The `\CameraUtil` class was developed as a mediator to simplify the interface to the Android `\Camera`. For example, it performs additional checks such as checking the orientation of the phone.

A DAO class called `\FaceDAO` was developed to mediate calls to the Android `\Javafaces` library. Images captured using `\CameraUtil` are validated using the DAO. The value returned from the `\Javafaces` library is the Euclidean distance between collected data and the registered owner. This distance is handled in exactly the same way as the voice recognition mechanism [5.2.2.3](#).

By default, the Android API does not easily allow for a Camera picture to be taken without any sort of notification to the user. Both a shutter sound and a visual preview display should be present. The shutter sound can easily be disabled by not providing a shutter callback function when calling the `Camera.takePicture()` method.

Disabling the user preview of the camera was more difficult to achieve. The solution used with this prototype was to exploit an Android feature that allows to render the preview in a `\SurfaceTexture` object. This satisfies the API's requirement to have a visual display preview for the camera, while the `\SurfaceTexture` itself does not need to be displayed on screen. Therefore a picture can be taken from a background service without any interruption to the user.

Another problem encountered by the face recognition service is data sizes. When the `\Javafaces` library performed its algorithm, the device was running out of memory. This caused the app to be closed by the Android OS. To fix this issue, Bitmaps collected from the camera are scaled to 50% before they are processed by the library.

5.2.2.4 Location analysis

The mechanism is based on gathering location data and using it to generate a probability that the owner is present. This is implemented in the `\LocationService` class that extends the `\AuthMechService` abstract class. Data is collected periodically by using the `\LocationManager` provided by the Android API.

A DAO object is used to mediate calls to the Android API and manage the existing owner configuration. It is implemented in the `\LocationDAO` class. It offers functionality for gathering and saving location updates. It is developed to use the most accurate data provider. The Android API offers the following sources of collecting `\Location` data:

- `GPS_PROVIDER`: Collects data from the GPS.
- `NETWORK_PROVIDER`: Collects data from cell tower and WiFi access points.
- `PASSIVE_PROVIDER`: Passively collects data from other applications which receive `\Location` updates.

External libraries were not used for the authentication process. We have developed a primitive location analysis algorithm in the `\LocationAnalyser` class. During the configuration stage, which is a process managed by `\LocationActivity`, location data is sampled every 5 minutes and saved in internal storage. After the process has ended, each time a `\Location` is sent for authentication it is compared with all the locations saved during the configuration process. The final result is the minimum Euclidean distance between the current `\Location` and any other saved `\Location`.

When the service is started by `\UAService`, its `\onCreate()` method spawns an authentication thread. This thread periodically requests the current location using the DAO. Data is returned in a `\Location` object and is provided as input to the `\LocationAnalyser`. The result of this operation is an Euclidean distance which gets converted to a percentage using a threshold, just as in the previous mechanisms. The value is stored by the service for future requests from `\UAService`.

Just as mentioned in section 5.2.2, the mechanism has a lower confidence level. Although being in a known location provides some confidence that the token was not stolen, it does not offer any guarantees that the device is still with its owner.

5.2.3 Owner configuration

There are a number of Activity components that are used in the configuration of the prototype. Each authentication mechanism has a corresponding Activity that can be started from the main `\PicoUserAuthenticator` Activity. These are used to register the owner biometrics based on which the mechanisms output their confidence levels.

Each configuration activity uses the same DAO class as the mechanism service. The DAO is used for storing the owner data once it was collected. Given that the overall size of the data is relatively small, the files are kept in internal storage. The Linux Android permissions mechanism guarantees its confidentiality and integrity, and therefore further encryption is not necessary.

5.3 Conclusion

We have described the Activity and Service components developed for the prototype, as well as their communication flow. We have ported two biometric libraries and developed a location analysis mechanism. DAO objects facilitate accessing owner configuration files and the interface with auxiliary mechanisms. An overview of the app design can be seen in figure ??.

A few limitations of the prototype is the lack of explicit authentication mechanisms. Another issue is the precision of the biometric mechanisms, in the lack of better libraries. However, due to the modular design of the application, existing mechanisms can be improved simply by importing a new library and modifying its DAO. The existing set of mechanisms can easily be increased by creating a new class that extends `\AuthMechService`", and developing the desired algorithm. In order to be managed by `\UAService`", the new mechanism needs to be included in the `\UserAuthenticator.initAvailableDevices()`" method.

5.4 Related work

Liang Cai et al [26] analyse ways of protecting users from mobile phone sensor sniffing attacks. The authors design a framework used for protecting sensor data from being leaked. From a security perspective the user should not to be trusted with granting permissions to different applications. An important point made in this paper is that malware may deny service to legitimate applications (such as our prototype) by creating a race condition for acquiring a sensor lock. The solution proposed by the authors would be an user notification, allowing for the owner to decide which application acquires the lock. A suggestion to this approach would be to allow for different priority levels, such that malware applications would not acquire the lock in a race condition, or even more, would lose it when a high priority application such as the Pico authenticator would require sensor data.

The paper by Derawi et al [22] presents the feasibility of implementing gait authentication on Android as an unobtrusive unlocking mechanism. According to the definition offered by the authors "gait recognition describes a biometric method which allows an automatic verification of the identity of a person by the way he walks". The Android implementation developed by the authors has an equal error rate (EER) of 20%. Dedicated devices have an EER of only 12.9%, and the main cause for this is the sampling rate available at that time (2010). They have used a Google G1 phone with approximately 40-50 samples per second. This is much inferior to dedicated accelerometers that sample data at 100 samples per second. However, by conducting personal experiments with the Accelerometer of a Google Nexus 5 phone, using the highest sampling setting (`SENSOR_DELAY_FASTEST`) the rates go above 100 samples per second. Therefore the current performance of the algorithm paper should be closer to 12.9%.

Ming et al [27] present in their paper how to improve speaker recognition accuracy on mobile devices in noisy conditions. This approach uses a model training technique based on which missing features may be used to identify noise. The focus of the paper is

designing and implementing a biometric mechanism, and is therefore outside the scope of this dissertation project.

Another technique in performing speaker recognition involves using voiceprints. These are a set of features extracted from the speaker sample data. Kersta [28] explains the mechanism in more detail. The benefit of having feature extraction based on a voice sample, as opposed to a different voice recognition mechanism, is that voiceprints do not require knowing any secrets. The speaker doesn't have to reproduce a voice sample. This increases the usability of the mechanism in scenarios required by the Pico authenticator. However, a downside to this approach is that it makes replay attacks easier to perform. Any recording of the user is sufficient for an attacker to trick the biometric mechanism.

A popular paper on face authentication was written by Turk and Pentland [29]. The biometric authentication process is based on the concept of eigenfaces. Eigenfaces are a name given for the eigenvectors which are used to characterise the features of a face. These features are projected onto the feature space. Using Euclidean distances in the feature space, classification can be performed to correctly identify individuals. An implementation of this mechanism was developed for the Pico unlocking scheme prototype.

An unconventional authenticating mechanism is presented by Clarke and Furnell [21]. They use keystroke analysis in order to make predictions regarding the user of the phone. This mechanism is unobtrusive and authenticates users during normal interactions such as typing a text message or a phone number. It is based on a neural network classifier, reporting an EER of 12.8%. Input data used for classification is composed out of timings between successive keystrokes, and the hold time of a pressed key.

Evaluation

5.5 Framework evaluation

We will continue by evaluating the new proposed scheme with the token unlocking framework defined in the previous chapter.

Memory-effortless: Satisfied

None of the authentication mechanisms require any sort of known secret. Authentication is granted based on biometrics and behavioural analysis.

Nothing-to-carry: Quasi-satisfied

This property is only quasi-satisfied due to the fact that it relies on the implementation of the design. Ideally all authentication data should be gathered from an unified device containing the Pico. Alternatively however, the scheme can be implemented using individual sensors which the owner would have to carry, which is why the property is not fully granted.

Easy-to-learn: Satisfied

In order to satisfy Pico's property of continuous authentication, all mechanisms part of the scheme I developed also need to have this property. Therefore the authentication process is non-transparent to the user, and therefore there is nothing to learn.

Efficient-to-use: Satisfied

The authentication data is collected either at fixed time intervals, or is triggered during special events. The authentication process however, does not fully depend on recent data. A response may be generated without any recent authentication data. Therefore the time spent by the mechanism to generate a response is immediate.

Infrequent-errors: Quasi-satisfied

Given that the scheme depends on biometric mechanisms, the quality of the errors is as good as the underlying biometrics. If the scheme cannot generate a high

enough confidence an explicit biometric challenge will be issued for the user to satisfy. Since the original biometric mechanisms do not have this property, to some extent neither will the scheme I have designed. However, the scheme is combining multiple biometrics results with different score weights based on importance and accuracy. This is much more likely to be accurate, which is why I will mark this as Quasi-satisfied. For a more accurate response, the design needs testing with a high quality prototype.

Easy-recovery-from-loss: Not-satisfied

Token based mechanisms in general do not have this property due to the inconvenience of replacing the token. In our case, the property is also not satisfied. The user would have to re-acquire a new token and reconfigure the owner's biometric data. Furthermore based on the mechanism, such as location settings or gait recognition, the token is likely to require an adaptation period.

Availability: Satisfied

Some mechanisms are not always available even though enabled, especially due to the continuous authentication property. As an example gait recognition while sitting in an office. However, the scheme may use a multitude of mechanisms with the unlikelihood that all of them are unavailable. For instance location history may predict with a certain confidence that the owner still in possession of the token. This property is aided by the explicit authentication mechanism which requires explicit input from the user.

Accessible: Satisfied

Due to the fact that the scheme is based on multiple biometrics and location settings, I consider this property to be Satisfied or as a very least Quasi-satisfied. The scheme functions based on available biometrics, without having any predefined solutions. It is highly unlikely that the owner cannot generate any of the available biometric inputs, especially for some such as "face recognition".

Negligible-cost-per-user: Quasi-satisfied

This property depends on the way in which the scheme is implemented. If the implementation is based on high quality sensors embedded in items of clothing and such, then the property is not satisfied. If the implementation reuses sensors that the user already possesses, then the property is fully satisfied as the cost is 0. An example of such an implementation would be an Android application/service possibly using the future Google Glass hardware.

Mature: Not satisfied

This property is not satisfied as the project is at the level of a work in progress

prototype. The design is quite fresh and was not implemented by any third party. Neither was it reviewed by the open source community or has had any user feedback.

Non-proprietary: Satisfied

Anyone can implement the scheme without any restrictions such as royalty checks or any other sort of payment to anyone else.

Resilient-to-physical-observation: Satisfied

Since the mechanism is based on biometric data, simple observations from an attacker cannot lead to compromising the user's authentication to the token. The attacker would have no way of reproducing the input through simple observation.

Resilient-to-targeted-impersonation: Quasi-satisfied

Saying that the scheme Quasi-satisfies this property is a bit generous. Each of the mechanisms is vulnerable to a replay attack. An attacker may record one of the user's biometric and replay it as a token input. However, given that the token uses multiple mechanisms, some of which being location based, this is a highly unlikely occurrence. The only vulnerable point would be the explicit authentication mechanisms, which carry a lot of weight.

Resilient-to-throttled-guessing: Satisfied

The amount of throttled guessing required for the user to break one of the biometric mechanisms is far too large for this to actually be a threat.

Resilient-to-unthrottled-guessing: Satisfied

Given that the Resilient-to-throttled-guessing property is satisfied, this property is also satisfied.

Resilient-to-internal-observation: Satisfied

This property does not apply to this scheme.

Unlinkable: Not-satisfied

Just as any of the biometric mechanisms, this property is not satisfied by the mechanism. The authentication data maps uniquely to the owner of the token.

Continuous-authentication: Satisfied

The mechanism was designed with continuous authentication in mind. Data is collected periodically with a confidence weight decaying over time. This allows for the token to be used at any time based on current existing data. The only exception breaking the model would be the explicit authentication mechanisms, but these could only be triggered at the beginning of an authentication process using the token.

Property	Picosiblings	Proposed scheme
Memory-efficient	Satisfied	Satisfied
Nothing-to-carry	Not-satisfied	Quasi-satisfied
Easy-to-learn	Satisfied	Satisfied
Efficient-to-use	Quasi-satisfied	Satisfied
Infrequent-errors	Quasi-satisfied	Quasi-satisfied
Easy-recovery-from-loss	Not-satisfied	Not-satisfied
Availability	Satisfied	Satisfied
Accessible	Not-satisfied	Satisfied
Negligible-cost-per-user	Not-satisfied	Quasi-satisfied
Mature	Not-satisfied	Not-satisfied
Non-proprietary	Satisfied	Satisfied
Resilient-to-physical-observation	Satisfied	Satisfied
Resilient-to-targeted-impersonation	Satisfied	Satisfied
Resilient-to-throttled-guessing	Satisfied	Satisfied
Resilient-to-unthrottled-guessing	Satisfied	Satisfied
Resilient-to-internal-observation	Satisfied	Satisfied
Unlinkable	Satisfied	Not-satisfied
Continuous-authentication	Satisfied	Satisfied
Multi-level-unlocking	Not-satisfied	Satisfied

Multi-level-unlocking: Satisfied

This property is fully satisfied by the authentication mechanism. It allows the token to grant access to different authentication accounts based on the precomputed level of confidence that the owner is present.

Let us continue by comparing the results of our proposed scheme with the original Picosiblings solution. The results are summarised in the following table. In the "Proposed scheme" column, properties which are highlighted in order to facilitate the comparison with the Picosiblings solution. The colour green means that the proposed scheme is better, red worse, and no colour means that both properties have the same value.

As the table shows, the proposed solution does not completely dominate the Picosiblings solution, and this is only because of the "Unlinkable" property. Given that our solution is fundamentally based on biometric data, this property could never be achieved. However, our solution performs better than Picosiblings in 5 other properties. Important points of improvement are accessibility, which makes the proposed scheme viable for a larger number of people. The Multi-level-unlocking property is another good improvement, allowing for an enhanced security model.

5.6 Conceptual design threat Model

An accurate threat model on the proposed unlock mechanism must start by analysing the set of assumptions made about the mechanism. From there we can identify available threats and how the scheme can be exploited in order to unlock the Pico without owner permission. Throughout the threat model we will explain how relaxing the initial set of assumptions may change the security outcome. Each model is analysed from an Availability, Integrity, and Confidentiality.

It is important to note that confidentiality is an important category in this evaluation. This is because the device will store sensitive biometric data which is directly linkable to the user. Losing this data, especially in plain-text, would disable the user from ever using the biometric device for which the data was leaked. This is due to the fact that the leaked data could always be replayed, successfully tricking the biometric mechanism.

In each subsection, the model will obviously only introduce issues with the mechanism. Therefore when reading a subsection, the issues are not only those currently presented, but also those from previous subsections that lead up to that point.

5.6.1 Dedicated device with dedicated sensors

We will start from the assumption that the unlock mechanism is integrated on the same device with the Pico. The device is assumed to be dedicated and runs no other software. Furthermore, the set of available sensors will also be integrated within the device. Alternatively there may also be peripheral sensors, with no way for an attacker to tamper with the communication to the authenticator.

Availability

From an availability point of view, an outside attacker cannot create a denial of service scenario. Interactions with the device are performed physically, so therefore the device cannot be made unavailable while in the possession of its owner. If the Pico would temporarily lose ownership, from a software perspective it would lock up due to mismatching biometric and location data, but would become available again in the presence of the owner.

Only hardware modification would affect data availability. Simply disconnecting the sensor would not affect the scheme's ability to generate viable results due to the fact that multiple biometrics are used. However an attacker could modify a sensor to output

wrong data, tricking it into saying the user is never the owner. This would create a successful denial of service attack path where a few sensors output that the owner is never present.

Integrity

Communication paths are not accessible from the outside and therefore cannot be tampered with in order to modify data. Furthermore the device is not running any other software and is therefore safe from any malware attacks.

Only physical tampering with the device would change data integrity. Modifying one of the sensor's and changing its output to some random data would be undetectable by the mechanism.

Confidentiality

No software access as well as no communication with the outside (i.e. wired communication) means that data is safe as long as the device is with its owner.

If the device were to be lost, Storage data should be kept encrypted, similar to the way Ironkey [] protects its data. Unfortunately an attack path may already be identified which is due to the fact that using this model the decryption key needs to be stored on the device. An attacker which has hardware access could therefore extract the key and decode the data. The original Picosiblings solution circumvented this approach by keeping

5.6.2 Dedicated device with shared sensors

We will relax the original set of assumptions by saying that the communication path with the sensors is no longer secure. Furthermore the sensors may be shared with other owners, via a wireless communication link for example. Another feasible scenario is that although sensors are located on the same device as the Pico, the Pico application is fully compartmentalised from the outside world.

What we are trying to stress with this scenario is that the sensors are no longer part of a trusted secure box, but are outside and communication with them, as well as their input may no longer be secure.

Availability

Since the sensors are no longer dedicated, other users may access sensor data. Depending on the hardware and software platform supporting the sensors, this may lead to a denial of service attack on the scheme. For example, if the sensors may only have one owner at a time, an attacker may request data from all sensors keeping them locked from the biometric authentication mechanisms. If the system is built in such a way, then there is nothing the scheme could do to prevent this other than keep the sensors constantly locked for itself. However since the model is built on the concept of shared sensors, this might not be a feasible solution.

Furthermore, communication paths are no longer dedicated. Whether the communication channel is radio or pure software, this introduces a new attack path. A "man in the middle" type of attack may be performed where information data from the sensors is dropped and replaced with bad data. This would create a scenario similar to the one in the previous section, but without the need for physically modifying the sensors.

Integrity

Having shared communication paths with the sensors means that data integrity may be compromised from outside. This goal would be achieved in the previous model only by physically modifying the sensors. Furthermore if the sensors are on the same device as the Pico, malware may modify output data leading to unsuccessful mechanism authentication.

Since Pico and the authenticating mechanism are fully compartmentalised from the outside, their communication is still secure. This compartmentalisation however needs to include all types of storage and communication.

Confidentiality

Unfortunately having shared sensors introduces quite a big confidentiality issue. Given that the sensor data required for authentication is shared, nothing would stop an attacker from collecting just as the Pico unlocking mechanism would. This data could then be replayed to the authenticator in order to unlock the Pico.

This is quite a critical issue. An example of feasible attack pattern would be. A piece of malware analyses when the sensors are locked, and makes assumptions as to when the

Pico authenticator is locking them. Based on these assumptions the malware then captures sensor data immediately after the lock was released therefore capturing a possibly valid sample of data.

A more elaborate piece of malware could detect patterns such as time intervals or events that trigger sensor locking. Knowing these patterns it could therefore lock the sensors and gather data just before the Pico authenticator would, and then trick the authenticator by sending it a replay or possibly modified data.

Yet another scenario in these circumstances would be to send the Pico authenticator constant bad data and anticipate the trigger of an explicit authentication request to the user. By locking the sensors at that key time the piece of malware could acquire a high quality data sample. Since most of the mechanisms used by the scheme are biometrics, that data sample would represent permanent damage to the user, as an authentication mechanism using that type of biometric could be replayed in any circumstance.

Since the Pico unlocking mechanism is fully compartmentalised, access its storage is secure and therefore any stored credentials are fully protected.

5.6.3 Insecure communication with Pico

This is a special case model which assumes that Pico and the authenticator we have developed are communicating over an insecure channel. The only element we need to consider is the communication between the two participants.

Availability

To do.

Integrity

To do.

Confidentiality

To do.

5.6.4 Shared device with shared components

We will relax the model even more in order to better fit reality constraints when implementing the mechanism. In this model, Pico and its authentication mechanism reside in a computing model with shared storage resources. The security of Pico and its authenticator may only be as good as the underlying OS. In order to have a meaningful use-case scenario.

Availability

To do.

Integrity

To do.

Confidentiality

To do.

5.6.5 Proposed secure implementation

A secure proposed implementation is viable using an Android telephone running a TrustZone enabled ARM processor available in ARMv6KZ [] and later models. This device would essentially be divided into two "worlds": the normal world running the untrusted Android OS, and the trusted world running a small operating system written for TrustZone. Both operating systems are booted at power up. In addition the TrustZone OS loads a public/private key pair which is inaccessible from Android.

Ideally Pico would be implemented with its authenticator within TrustZone. This would essentially guarantee complete separation from a memory perspective leaving any sort of malware attack impossible via memory.

Persistent memory is however required in order to store data for each individual biometric mechanism used in the authentication scheme. Unfortunately this type of memory is not protected by the TrustZone OS and constitutes a way for a third party to attack the scheme. However, we could use the TrustZone OS key pair in order to encrypt biometric data on disk. Even though this data is available from Android it would be fully

confidential. If properly stored within Android, the OS may even protect its integrity from outside attacks.

Let us consider however that the Android OS has been completely compromised by the attacker and is therefore "hostile". Under these circumstances data confidentiality can still be fully guaranteed. The TrustZone public key could still be used in order to encrypt the biometric data before writing it to disk. Attacks from a memory perspective may only be performed by modifying data stored on disk. This may only lead to a denial of service for the owner, but not a confidentiality breach.

Let us briefly discuss any issues using the availability-integrity-confidentiality framework.

Availability

Only plausible attacks are denial of service through deleting biometric cache files from disk. This would require constant reconstruction for the Pico scheme, making the Pico unavailable.

Integrity

Data integrity may only be altered from cache files on disk.

Confidentiality

No known attacks on data confidentiality other than capturing sensor data just as the authenticator would. However this would be possible with or without the Pico being present.

5.7 Threat model

Even though the scheme implementation is a proof of concept, we will continue by analysing different threat models. This will reveal any flaws behind the concept, allowing for a more robust future implementation.

The purpose of the Pico token is to provide a robust authentication mechanism, without the use of any secrets for the owner to remember. Where the Pico unlocking scheme is, is correctly identifying the owner of the Pico. Attacks may be performed in the form of malware installed on the device while still within the possession of its owner. The main threat however comes from an attacker having physical access to the user's Pico.

It is important to note that since this is a purely software implementation, physical access may mean either that the attacker is in possession of the phone, or that it may replicate the secrets of the victim's Pico on a separate device. Replicating the Pico secrets would clearly create much more damage for the user from a cost perspective. A total reset of authentication credentials would be necessary for all accounts registered with the Pico device.

5.7.1 Prototype threat model

Let us now continue by studying the threat model of the Pico authenticator application. We will consider the security mechanisms presented above as the predefined assumptions made in this model. In order to reduce the threat space we will consider the application is running on a hand held device running Android 4.4.2 with all recent updates.

Availability

Breaking the scheme's availability if the device is in the possession of the attacker is relatively trivial. The application can be uninstalled, or the application data cache can be cleared, therefore removing the owner biometric models for the different mechanisms. Furthermore, in this case the owner is already no longer in possession of their Pico, so basically the Pico is already made unavailable..

Let us continue however and study what can be achieved from a DoS perspective by the attacker from the perspective of the individual user app accounts, which would need to be reset by the owner. In order to gain any sort of access and make credentials reset not possible, or at least have a chance in doing so, the attacker would have to unlock the Pico.

From a malware attack perspective data used by the authentication should not be modifiable. This would guarantee that all mechanisms have their cached biometric data available at all times and may function properly. Due to the Linux permissions mechanism and the fact that each application has its own UID and GID, data stored in internal memory should not be readable or writeable by any other user level application in the system. If however the device is rooted and the owner is misled into granting root privileges to another application, then the security model would be broken and the data would be exposed. This could lead to deletion which would make the mechanisms not function properly, resulting in a DoS attack.

Integrity

Just as mentioned in the Availability section, the authenticator should be safe against any data accesses from other applications as long as the application does not have root privileges. This would allow malware to break the integrity property of the data.

From a data flow point of view Intents used for communication within the authenticator as well as with the Pico application are not modifiable. Furthermore Intents are not broadcasted using the implicit Android broadcast mechanism, which makes them impossible to replay or even intercept.

Confidentiality

Considering a circumstance where a root malware process would have access to the authenticator's data stored on disk, this would not lead to a direct compromise of the owner biometric data. All cached files are stored in internal memory in encrypted format. The mechanism used for encryption is RSA, with the private key stored using the Android KeyChain API.

On a rooted device however, the encryption layer provides only another bi-passable layer of security. With root access, an application could retrieve the master key of the KeyChain and use it to retrieve or private key and decode the owner's biometrics.

From a data flow perspective, internally the Pico authenticator uses an self-developed broadcast system. Client processes need to register with the broadcaster, such as the UAService, in order to receive updates. This ensures data confidentiality throughout the system. Furthermore, the authenticator and Pico should be released under the same author. This would allow locking the application from outside Intents as well as interaction with different components. Sandboxing communication is always a desirable property from a confidentiality perspective.

The paper by Adrienne Porter Felt et al [30] shows that according to their surveys only 17% of users pay attention to the Android permissions dialogue, and only 3% understand what each permission represents. A malware application which has granted full permissions gets pass the Bouncer and is installed as an application. Even so, due to the Linux permission model adopted by Android, the confidentiality of the authenticator's data would not be compromised. Instead however, the malware application may collect all relevant on its own from the user, allowing for a powerful replay attack in the future.

5.8 Future work

The application was implemented as a proof of concept. It is developed in order to show that different data may be obtained without the owner's knowledge. Additional improvements can be made in order to increase the confidence level of the authenticator. Furthermore, due to time constraints and unavailability of free to use biometric libraries, a number of mechanisms were not implemented. The list can easily be extended by simply creating a class which extends the "AuthMechService" abstract class.

One way to improve the voice recognition mechanism would be to start sampling data whenever a call is active. This would increase the chances of capturing an accurate sample of the owner's voice. In this context, a better voice recognition library can be used, which supports multiple speakers and/or ignores background noise. If such a library is not available, we can rely on the fact that most of the times people take turns when speaking. For the duration of the call, with a high enough sampling frequency, the individual sampling voice of both participants should be captured. However, it is important to take into account a situation in which the thief is calling the owner on a different phone in order to unlock his Pico.

Immediate improvements can be made to the face recognition mechanism. Just as recommended in the description of the mechanism's implementation, another library which provides more meaningful face coordinates may be used for face detection. Alternatively, and preferably, a different library which performs both face detection and recognition can be integrated with the mechanism.

Another improvement for the face recognition mechanism would be from the data sampling perspective. Instead of capturing images at a fixed interval, pictures should be taken only when the phone unlock event is triggered. While the phone is unlocked it is highly likely that the user will face its front camera. This would provide better chances of processing meaningful data.

Appendix Title Here

Write your Appendix content here.

Bibliography

- [1] Anne Adams and Martina Angela Sasse. Users are not the enemy. *Communications of the ACM*, 42(12):40{46, 1999.
- [2] Robert Morris and Ken Thompson. Password security: A case history. *Communications of the ACM*, 22(11):594{597, 1979.
- [3] Philippe Oechslin. Making a faster cryptanalytic time-memory trade-o . In *Advances in Cryptology-CRYPTO 2003*, pages 617{630. Springer, 2003.
- [4] Je Jianxin Yan, Alan F Blackwell, Ross J Anderson, and Alasdair Grant. Password memorability and security: Empirical results. *IEEE Security & privacy*, 2(5):25{31, 2004.
- [5] Joseph Bonneau, Cormac Herley, Paul C Van Oorschot, and Frank Stajano. The quest to replace passwords: A framework for comparative evaluation of web authentication schemes. In *Security and Privacy (SP), 2012 IEEE Symposium on*, pages 553{567. IEEE, 2012.
- [6] Nathan L Clarke, Steven M Furnell, Phihp M Rodwell, and Paul L. Reynolds. Acceptance of subscriber authentication methods for mobile telephony devices. *Computers & Security*, 21(3):220{228, 2002.
- [7] Frank Stajano. Pico: No more passwords! In *Security Protocols XIX*, pages 49{81. Springer, 2011.
- [8] Oliver Stannard and Frank Stajano. Am i in good company? a privacy-protecting protocol for cooperating ubiquitous computing devices. In *Security Protocols XX*, pages 223{230. Springer, 2012.
- [9] Dinei Florencio and Cormac Herley. A large-scale study of web password habits. In *Proceedings of the 16th international conference on World Wide Web*, pages 657{666. ACM, 2007.
- [10] Adi Shamir. How to share a secret. *Communications of the ACM*, 22(11):612{613, 1979.

- [11] Frank Stajano. The resurrecting duckling. In *Security Protocols*, pages 183{194. Springer, 2000.
- [12] Hugh Wimberly and Lorie M Liebrock. Using fingerprint authentication to reduce system security: An empirical study. In *Security and Privacy (SP), 2011 IEEE Symposium on*, pages 32{46. IEEE, 2011.
- [13] Nathan L Clarke and Steven M Furnell. Authentication of users on mobile telephones{a survey of attitudes and practices. *Computers & Security*, 24(7):519{527, 2005.
- [14] Nathan L Clarke, SM Furnell, Benn M. Lines, and Paul L Reynolds. Using keystroke analysis as a mechanism for subscriber authentication on mobile handsets. In *Security and Privacy in the Age of Uncertainty*, pages 97{108. Springer, 2003.
- [15] Gregory D Williamson and GE Money-America's. *Enhanced authentication in on-line banking*. PhD thesis, Utica College, 2006.
- [16] Elena Vildjiounaite, S-M Makela, Mikko Lindholm, Vesa Kyllonen, and Heikki Ailisto. Increasing security of mobile devices by decreasing user effort in verification. In *Systems and Networks Communications, 2007. ICSNC 2007. Second International Conference on*, pages 80{80. IEEE, 2007.
- [17] William Enck, Machigar Ongtang, Patrick Drew McDaniel, et al. Understanding android security. *IEEE Security & Privacy*, 7(1):50{57, 2009.
- [18] Steffen Liebergeld and Matthias Lange. Android security, pitfalls and lessons learned. In *Information Sciences and Systems 2013*, pages 409{417. Springer, 2013.
- [19] Patrick Gage Kelley, Sunny Consolvo, Lorrie Faith Cranor, Jaeyeon Jung, Norman Sadeh, and David Wetherall. A conundrum of permissions: Installing applications on an android smartphone. In *Financial Cryptography and Data Security*, pages 68{79. Springer, 2012.
- [20] NJ Percoco and S Schulte. Adventures in bouncerland: Failures of automated malware detection within mobile application markets. *Black Hat USA 2012*, 2012.
- [21] Nathan L Clarke and SM Furnell. Authenticating mobile phone users using keystroke analysis. *International Journal of Information Security*, 6(1):1{14, 2007.
- [22] Mohammad Omar Derawi, Claudia Nickel, Patrick Bours, and Christoph Busch. Unobtrusive user-authentication on mobile phones using biometric gait recognition. In *Intelligent Information Hiding and Multimedia Signal Processing (IIH-MSP), 2010 Sixth International Conference on*, pages 306{311. IEEE, 2010.

- [23] Mark Burge and Wilhelm Burger. Ear biometrics. In *Biometrics*, pages 273{285. Springer, 1996.
- [24] Zhichun Mu, Li Yuan, Zhengguang Xu, Dechun Xi, and Shuai Qi. Shape and structural feature based ear recognition. In *Advances in biometric person authentication*, pages 663{670. Springer, 2005.
- [25] Maycel Isaac Faraj and Josef Bigun. Motion features from lip movement for person authentication. In *Pattern Recognition, 2006. ICPR 2006. 18th International Conference on*, volume 3, pages 1059{1062. IEEE, 2006.
- [26] Liang Cai, Sridhar Machiraju, and Hao Chen. Defending against sensor-sni ng attacks on mobile phones. In *Proceedings of the 1st ACM workshop on Networking, systems, and applications for mobile handhelds*, pages 31{36. ACM, 2009.
- [27] Ji Ming, Timothy J Hazen, James R Glass, and Douglas A Reynolds. Robust speaker recognition in noisy conditions. *Audio, Speech, and Language Processing, IEEE Transactions on*, 15(5):1711{1723, 2007.
- [28] Lawrence George Kersta. Voiceprint identi cation. *The Journal of the Acoustical Society of America*, 34(5):725{725, 2005.
- [29] Matthew A Turk and Alex P Pentland. Face recognition using eigenfaces. In *Computer Vision and Pattern Recognition, 1991. Proceedings CVPR'91., IEEE Computer Society Conference on*, pages 586{591. IEEE, 1991.
- [30] Adrienne Porter Felt, Elizabeth Ha, Serge Egelman, Ariel Haney, Erika Chin, and David Wagner. Android permissions: User attention, comprehension, and behavior. In *Proceedings of the Eighth Symposium on Usable Privacy and Security*, page 3. ACM, 2012.